The Agentic Organisation

The Agentic Organisation

A complete AI transformation framework for European CEO's

© 2026 The Agentic Organisation. All rights reserved.

This work is for distribution via Amazon KDP and IngramSpark.

Contents

Introduction	6
Chapter 1 - The Largest Shift Since the Industrial Revolution	11
Chapter 2 - Leadership Is Not Optional	26
Chapter 3 - What the Agentic Organisation Looks Like	40
Chapter 4 - The Four Tiers of Transformation	55
Chapter 5 - The AI Transformation Framework	72
Chapter 6 - The Operational Layer	81
Chapter 7 - Foundations of the AI Transformation	92
Chapter 8 - Value, Capabilities, Trust	105
Chapter 9 - Europe's Distinctive Edge	120

Introduction

The question most boardrooms are asking is the wrong one. "How do we adopt AI?" sounds strategic. It leads to pilots, proofs of concept, and a lot of activity. It also leads, in the vast majority of organisations, to no sustained impact.

The right question is harder. "How do we lead an organisation that can absorb and act on AI?" That question does not belong in the IT department. It belongs in the room where strategy, governance, and culture are decided. This book is written for the leaders in that room.

I have spent over twenty-five years helping organisations create value from technology, long before AI became a boardroom imperative. The pattern I see now is consistent. The technology works. The strategies are often sound. The failure is in the organisational response. European CEOs are not underprepared because they lack ambition or intelligence. They are underprepared because the mental model most of them are applying, incremental and project-based, is structurally incapable of capturing what this moment demands.

This is the largest paradigm shift since the industrial revolution. Treating it as another software rollout is the single most expensive mistake a leadership team can make.

Why a European perspective

This book is written from and for Europe deliberately. The European context is not a constraint to work around, it is a different playing field with distinct advantages. The EU AI Act gives one harmonised framework where other markets face fragmented or absent rules; public trust in the EU to regulate AI is higher than in the US or China, and evidence shows that well-designed regulation can unlock demand rather than block it. European

labour dialogue—works councils, co-determination, collective agreements—shapes how AI enters the workplace in ways that build workforce acceptance and, as some leading firms have found, surface concerns that customers share. GDPR and years of data-governance discipline have created a head start on the data foundations that AI transformation requires. None of that makes the transformation easy, but it means the playbook that works in Europe is not the same as the one written for other regions.

Who this book is for

This book is for CEOs, C-suite executives, and board members of organisations in Europe. It is for leaders who want strategic clarity, not hype. Who understand that the real work of AI transformation happens in the organisation, not the algorithm. Who are willing to ask the harder questions and to lead from the top, because the evidence is unambiguous: transformation that is delegated fails; transformation that is CEO-led succeeds at a multiple of the rate of the rest.

What you will find here

The central idea is the *Agentic Organisation*: an organisation redesigned around human–AI collaboration, where humans and AI agents work as a unified system rather than tools bolted onto legacy processes. Getting there is not a technology project. It is a leadership project with a clear sequence: understand the shift, own the leadership requirement, learn what the destination looks like, then build it in stages with the right foundations and the right measures.

Part I lays out the burning platform. Why this moment is different. Why leadership cannot be optional. Part II defines the model: what the Agentic Organisation looks like, how transformation moves from individual to team to department to organisation, and how to decide where to invest using the AI Decision Framework (value, capabilities, and trust). Part III turns to Europe: how regulation, trust, and pragmatism become advantages, and

what leaders in six European companies actually did when the stakes were high. Part IV is the playbook: how to build your Agentic Organisation from foundation through activation to scale, what must be in place before and alongside transformation, what to measure, and why the final chapter is not a summary but a decision.

You will not find a catalogue of tools or a vendor guide. You will find frameworks you can use in the room, evidence from research and from the field, and a practical sequence with clear milestones. The goal is to give you the mental model and the next steps so that you can lead, not delegate.

When this book was written

This book was completed in March 2026. That matters. The data, studies, and market evidence cited here reflect what was known and interpretable at that moment. AI adoption curves, regulatory implementation, and the pace of agentic deployment are all moving quickly; by the time you read this, some numbers will be outdated, some expectations will have been overtaken by events, and new uncertainties will have emerged. The analysis and recommendations are based on today's interpretation—with many uncertainties acknowledged. Use the frameworks and the sequence; treat the specific figures and timelines as a snapshot, not a guarantee. The mental model and the leadership imperative, however, are built to last beyond any single data point.

Practicing What We Preach

Writing a book about the agentic organisation while relying entirely on traditional methods would have been, at best, inconsistent. At worst, it would have undermined the very argument this book makes. So I made a deliberate choice early on: I would use AI as a genuine working partner throughout the research and writing process, not as a shortcut, but as a demonstration of what human-AI collaboration actually looks like in practice.

In concrete terms, that meant using AI to synthesize research across hundreds of sources, challenge the logic of emerging frameworks, stress-test arguments, and help translate complex ideas into clear prose. Every framework in this book originated from my own thinking, built on over 25 years of working with organizations navigating technology-driven transformation. But AI helped me think faster, go deeper, and pressure-test conclusions in ways that would have taken far longer working alone. The result is a book that is genuinely better for it.

I want to be transparent about what this means for you as a reader. The ideas, the judgments, the recommendations, and the responsibility for what is written here are mine. AI did not determine what matters or what to conclude. What it did was act as a tireless research partner, a rigorous editor, and a sounding board that never got tired of the question "but is that actually true?" If you recognize some of that dynamic in how your own teams might one day work alongside AI agents, that is entirely intentional.

In that sense, this book is not just about the agentic organisation. It is also a small proof of concept for it.

How to use this book

Read it in order if you want the full argument. Use the table of contents if you want to jump ahead: go to the AI Transformation Framework (Chapter 5) or the implementation playbook (Chapter 11) if you need a specific tool. Each chapter ends with reflection questions you can use with your team or your board. The 9 lessons in Chapter 10 (drawn from case studies) show how leaders decided under pressure; the lessons are meant to transfer, not the cases to be copied literally. Your context will differ, but the underlying patterns do transfer.

The cost of inaction is not zero. It is already accumulating in cost structure, talent pipeline, and competitive positioning. The organisations that will matter in three years are the ones that treat this as a leadership decision and act accordingly. This book is written to help you be one of

them.

Chapter One

The Largest Shift Since the Industrial Revolution

The Ninety-Five Per Cent Paradox

Ninety-five per cent of enterprise AI implementations fail to deliver sustained impact. Not because the technology does not work. Not because the strategy was wrong. Because organisations are treating the largest paradigm shift since the industrial revolution as though it were another software rollout.

The data on this is unambiguous. The MIT NANDA report, drawing on over 300 AI initiatives across 52 organisations and 153 senior leader surveys, found that 95% of enterprise generative AI implementations fail to deliver sustained, measurable impact. A global survey of more than 1,000 senior executives found that only 5% of companies achieve AI value at scale, whilst 60% achieve no material value whatsoever.

RAND puts the AI project failure rate at over 80%, double the rate of non-AI IT projects. S&P Global reports that 42% of companies scrapped most AI initiatives in 2025, up from 17% the year before.

These are not preliminary findings from a single study. They are converging conclusions from independent research programmes spanning thousands of organisations. And yet, eight in ten companies already use generative AI. Equally many report no significant bottom-line impact. The tools are everywhere. The value is not.

The paradox is striking because the technology does work. The companies achieving value at scale are achieving it dramatically: 2-3x conversion improvements, 80% reductions in processing time, hundreds of millions in measurable value. A Singapore bank grew its AI economic value from S$180 million to S$750 million in two years. A Dutch bank achieved a 90% pilot-to-production rate against an industry average of 30%. A German insurer reduced claim processing time by 80% through a seven-agent AI workflow launched in under 100 days.

These outcomes are real and replicable. The question is not whether AI can produce results of that order. It is why most organisations cannot access them despite considerable effort and investment.

The answer is one of the oldest problems in business: applying the wrong mental model to a categorically different kind of change. Most organisations are doing what early automobile buyers did when they called it a "horseless carriage": defining the new technology in terms of the old one. They are putting AI on top of existing processes, existing org charts, existing decision-making structures, and asking it to make things faster. That produces faster versions of the same thing. It does not produce transformation.

McKinsey, in its landmark September 2025 analysis of agentic AI, described this moment as "the largest organisational paradigm shift since the industrial and digital revolutions". That framing is not consulting hyperbole. It is an analytical claim with precise implications.

Paradigm shifts do not yield to incremental responses. The organisations that thrived during the industrial revolution did not optimise their

horse-drawn carriages. They redesigned their operations around the principles that steam and mechanisation made possible. Those that treated the new technology as an upgrade gained temporary efficiency, then fell behind organisations that had redesigned entirely. The same pattern is now unfolding with AI, at a pace that makes the industrial revolution look glacial.

AI task completion capability is doubling every four months. The window for first-mover advantage is not widening. It is narrowing. The question this chapter asks is not whether you believe that. The data settles it. The question is whether you are leading as though you do.

Why This Moment Is Different

Every technology wave produces bold claims about transformation. What makes this one categorically different is that the capability curve is exponential, the rate of acceleration is itself increasing, and the gap between leaders and laggards is already compounding in ways that are measurable and, for the laggards, increasingly difficult to reverse.

The numbers tell the story. Since 2019, AI task completion ability has doubled every seven months. Since 2024, the rate has accelerated to every four months. As of late 2025, AI systems can complete approximately two hours of autonomous work without supervision. The same trajectory, if it holds through 2027, puts that figure at four days: the equivalent of a mid-tenure employee operating independently for a full working week.

Each delay compounds. More compute, greater model efficiency, and fundamental changes in how AI systems are designed and deployed are driving what some researchers project could be a 10,000-fold increase in model capability within three years. The model that costs a department relatively little today does not simply become twice as capable in two years. It becomes capable of tasks that were impossible when the budget was

approved.

This is the critical distinction: the rate of improvement is itself accelerating. Organisations planning on a linear trajectory are building for a world that will not exist by the time their plans mature. The strategy approved in January may be structurally obsolete by September, not because the strategy was bad, but because the underlying capabilities shifted faster than the planning cycle could accommodate. For leaders accustomed to three-year technology roadmaps, this compression is disorienting. It is also the new reality.

> **Agentic AI**
>
> Agentic AI refers to AI systems that act autonomously to complete multi-step tasks, make decisions, and use tools or other agents to achieve a defined goal. Unlike generative AI, which produces outputs in response to individual prompts, agentic AI plans, reasons, and takes sequences of actions without continuous human instruction. Predictive AI handles logic; generative AI handles creativity; agentic AI exercises what BCG describes as the "frontal cortex" function: planning, judgment, and orchestration.

The organisations capturing this value are doing so visibly. Research from a global survey of more than 1,000 senior executives found that the top 5% of companies, those furthest ahead in AI, achieve 3.6 times the three-year total shareholder return of their peers, 1.7 times the revenue growth, and 1.6 times the EBIT margin. These leaders plan to spend more than twice as much on AI as the organisations trailing them.

Agentic AI already accounts for 17% of total AI value in 2025, with projections reaching 29% by 2028. The gap is not a snapshot. It is a trend line with compounding momentum.

> AI agent
>
> An AI agent is a software system that perceives its environment, plans actions, and executes tasks autonomously to achieve a defined goal. Agents can use tools, access data, communicate with other agents, and take action without step-by-step human instruction. In an organisational context, agents take on work that previously required human attention at every step: research, scheduling, analysis, compliance checking, decision preparation, and process execution.

The industry has moved through distinct phases, and each one has raised the stakes. In 2023, organisations experimented with generative AI on modest budgets and with considerable curiosity. In 2024, the "thousand flowers" phase arrived: every department launched its own AI initiative, chatbots proliferated, and individual productivity gains were real. But at the enterprise level, almost none of it compounded. Now, in 2025 and beyond, a small group of companies is focusing on the high-return transformations that produce structural change, whilst the rest remain scattered across initiatives that feel productive but do not add up.

For European organisations, the context carries additional weight. Only 20% of EU enterprises use AI today. Large enterprises sit at 55%, but the spread is dramatic: Denmark at 42%, Romania at 5.2%. Over 60% of European firms remain at the earliest stages of AI maturity.

The investment gap is real. US private AI investment in 2024 reached $109 billion. Europe produces only three of the world's top AI models, compared to 40 in the United States. But the picture is not one-sided. A 2026 EU Commission-backed report found that Europe allocates a higher share of venture capital to agentic AI than the US, with Paris, Berlin, Amsterdam, Stockholm, Munich, and Barcelona as leading hubs. The raw materials for European leadership exist. The question is whether they will be assembled into something that compounds.

The companies generating breakthrough value, wherever they sit, share a characteristic that has little to do with which AI vendor they chose or how large their technology budget was. What they share is a commitment to treat AI as an organisational transformation, not a technology deployment. The common trait across the top performers is leadership-led transformation: leaders who understood that AI required a fundamentally different kind of response, and who committed to providing it.

Against this, the majority pattern is familiar. Grassroots AI experimentation creates cultural momentum and generates genuine individual gains. But it does not self-organise into enterprise-wide impact. Without clear direction from the top, experimentation remains fragmented, siloed, and ultimately shallow.

The numbers make this gap vivid. Research confirms that 91% of employees already use generative AI at work. Only 1% of companies report AI maturity. Individual adoption massively outpaces organisational readiness, because individual adoption follows tools and convenience, whilst organisational maturity requires deliberate architecture.

The employee using AI to draft emails faster is productive. A thousand employees doing the same, in isolation, without shared infrastructure or coordinated process redesign, is not a transformation. It is a thousand separate improvements that never connect.

Industry research names this the micro-productivity trap: grassroots experimentation that sparks innovation and cultural momentum but does not, and cannot, self-organise into enterprise-wide impact. Without strategic direction from leadership, those gains remain fragmented and shallow, never aggregating into the kind of structural change that moves the numbers on a board report.

BCG's 10/20/70 rule makes the structural issue plain. Algorithms account for 10% of what is required to generate real AI value. The technology backbone is another 20%. The remaining 70% is people and processes: the

organisational decisions about how work is designed, how authority is allocated, how humans and AI agents collaborate, and how the organisation learns from its deployments. Most companies invest heavily in the 10% and treat the 70% as a downstream consequence. It is not. It is the primary variable.

This is the horseless-carriage problem in its purest form. Organisations are putting AI on top of existing processes, existing decision flows, and existing ways of working, then wondering why the results plateau. The technology is capable of something categorically different: not faster versions of current work, but fundamentally different ways of making decisions, flowing value, and collaborating across human and machine capability. Accessing that requires redesigning how the organisation works, not adding tools to how it already works.

One researcher captured the underlying challenge precisely: "Until we make an exponential shift in our mental model from spreadsheets to algorithms, AI tech will only give us an incrementally better spreadsheet." The shift required is not in the tools. It is in how leaders conceive of what the tools make possible.

Eric Kutcher, McKinsey's North America Chair, put the misframing plainly: "This is probably the biggest, most complex business transformation, but it's 80% business transformation and 20% tech transformation. That's different from how most people have thought about it".

Most have thought about it as the 20%. That is why 95% are failing.

The Governance Gap

The 95% failure rate is not primarily a technology problem. It is a governance and leadership problem.

Boards are failing to keep pace with the organisations they oversee. A Deloitte study of 695 board members across 56 countries found that 66% report limited to no knowledge or experience with AI: an improvement from 79% the prior year, but still a clear majority. Only 15% of boards receive AI-related metrics. Fewer than 25% have board-approved AI policies. Only 14% discuss AI at every board meeting. Meanwhile, 88% of their organisations already use AI in at least one business function.

The gap between what is happening inside the organisation and what the board understands is itself a strategic risk. This is not abstract. When a board cannot distinguish between a productive AI pilot and a structurally limited one, it cannot allocate capital, attention, or governance effectively.

When AI is discussed once per quarter rather than at every meeting, the pace of the technology outruns the pace of oversight. And when 66% of board members lack the knowledge to ask the right questions, the organisation's AI ambitions are governed by people who cannot evaluate what they are governing.

> ### The Agentic Organisation
>
> The Agentic Organisation is a company redesigned so that humans and AI agents work together as a unified system. It is not an organisation that uses AI tools. It is one whose operating architecture, decision-making processes, and workflow design are built for human-agent collaboration at scale. The Agentic Organisation does not add AI to the way it works. It works in a fundamentally different way because of AI.

The consequences of inaction are measurable. CEO-led AI transformations succeed at six times the rate of those delegated to IT or middle management. C-level executives deeply engaged with AI are twelve times more likely to be among the top 5% of companies winning with AI. The strongest single predictor of AI success is not the quality of the model,

the size of the budget, or the sophistication of the strategy. It is whether senior leadership is personally, visibly, and irreversibly committed to leading the transformation. Delegation is not a governance posture. It is a failure mode.

The board's role is not to become technical. Governance at the right level of abstraction means determining the organisation's AI posture, demanding specific metrics at the board table, and actively challenging whether the organisation is pursuing genuine transformation or a portfolio of disconnected pilots.

What does that look like in practice? Boards should track ROI by business unit, the percentage of processes that are AI-enabled, workforce reskilling progress, regulatory alignment, and human override rates in agentic systems. They should know how many AI initiatives have moved from pilot to production, and why the ones that have not are stalled.

They should be able to distinguish between activity and impact. These are not IT metrics. They are organisational health metrics, and they belong on every board agenda.

Three out of four executives rank AI as a top-three strategic priority, yet only a quarter report meaningful value from it. That gap is a governance failure, not a technology failure. When the board treats AI as a technology investment delegated to the CIO, it gets technology-investment outcomes: isolated pilots, modest efficiency gains, and no compounding effect across the organisation. When it treats AI as the strategic inflection point it is, the entire governance posture changes.

For European boards, the compliance dimension adds urgency that cannot be deferred. Full obligations under the EU AI Act for high-risk AI systems apply from 2 August 2026, with penalties of up to 35 million euros or 7% of global annual turnover for prohibited practices. Prohibitions on unacceptable-risk AI have already applied since February 2025. AI literacy training is already required. Any firm using AI in or affecting EU users,

regardless of where it is incorporated, falls within scope.

Eurostat data from December 2025 shows that 20% of EU enterprises with ten or more employees now use AI, up from 13.5% the prior year. Large enterprise adoption stands at 55%. But adoption without governance is exposure. The World Economic Forum has identified Europe's AI adoption lag as an urgent structural challenge, not a cyclical one. The compliance clock is running in parallel with the competitive clock. Boards that treat the EU AI Act as a legal department matter, rather than a governance imperative, are misallocating both risk and attention.

The board's first task is not to approve an AI strategy. It is to treat this for what it is: a strategic inflection point that demands the urgency, oversight, and accountability that characterisation requires. Every quarter of delay is a quarter in which the organisations that have already committed are compounding their advantage, and a quarter in which the governance gap between what the board understands and what the organisation needs becomes harder to close.

The Window Is Narrowing

This is not the cloud migration, the digital transformation, or the mobile-first pivot revisited at larger scale. It is a restructuring of how organisations create value, make decisions, and deploy human capability. The organisations already moving are compounding their advantages every quarter: in productivity, in talent, in data, in competitive positioning. The cost of waiting is not zero. It is already accumulating.

The leaders are not standing still. They are accelerating. Every quarter of compounding advantage in AI capability, data assets, and organisational learning widens the gap further. Research shows that the top 5% of companies already generate 2.7 times greater ROI from their AI investments, and they plan to invest more than twice as much again. This is

not convergence. It is divergence, and it is structural.

The organisations that will define their industries over the next decade are making irreversible commitments now, not waiting for clarity that will not arrive on schedule. For European leaders specifically, the window carries a dual character: the competitive pressure to act, and the regulatory infrastructure that, properly understood, provides a foundation for trusted AI at scale. Both clocks are running. Neither pauses for deliberation.

Recognising the scale of the shift is the first step. The harder question is whether the organisation has the leadership to match it. Who is accountable? Where does responsibility sit? What happens when that accountability is diffuse, delegated, or absent?

The next chapter takes that question directly.

Questions for the Board

> What would have to be true for your organisation to be in the five per cent achieving AI value at scale, and how much of it is true today?
>
> If AI task completion capability doubles again in four months, which of your current strategic assumptions become invalid?
>
> Ninety per cent of transformative AI use cases are stuck in pilot mode. How many of yours are? What is actually preventing them from scaling?
>
> When did your board last discuss AI as an organisational transformation rather than a technology investment? What was decided, and what action followed?
>
> CEO-led transformations succeed at six times the rate of those delegated downward. Who in your organisation is unambiguously accountable for this transformation?

Case Study: Siemens, The 177-Year-Old Company That Chose Transformation

In 2023, Siemens faced a choice that every large industrial organisation eventually confronts: what kind of AI company do you want to become? The first path was familiar: embed AI tools into existing manufacturing processes, generate efficiency gains, show progress to the board. This is the path most industrial companies have taken, and it is defensible. Siemens had the process complexity, the legacy systems, and the deeply embedded engineering culture to justify a cautious approach. The second path was categorically different: treat AI not as an optimisation layer but as the operating infrastructure of a redesigned industrial enterprise. This path was slower to justify, harder to govern, and required a level of commitment that most 177-year-old institutions are not built to make. Siemens chose the second path.

The industrial manufacturing sector faces a particular version of the AI adoption problem. Complex machinery, multi-decade equipment cycles, and safety-critical processes create genuine reasons to layer new technology cautiously onto proven systems. The argument for incremental AI adoption in industrial settings is not irrational: it is, in fact, the responsible-sounding default. Most of Siemens' peers chose it.

What made Siemens' context different was a recognition that the expertise gap inside industrial AI was not a technology problem. The barrier was not that AI lacked capability for industrial applications. The barrier was that the interface between AI systems and human operators, engineers who had spent careers developing proprietary knowledge about complex machinery, was broken. Line workers could not query AI systems in natural language. AI outputs were not embedded in the workflows where decisions were made. And AI models trained on generic data had no understanding of the specific physics, tolerances, and real-world constraints of Siemens industrial environments.

This diagnosis led directly to the design of Siemens' response. Rather than deploying standard AI tools into existing processes, Siemens committed to building industrial-grade AI infrastructure: models built to understand machines, workflows, and real-world constraints, with human interaction embedded at the point of work. The decision was not made once; it was made at three distinct levels simultaneously.

Committing to deep transformation over incremental AI adoption carried real costs. Building industrial-grade AI foundation models and a Digital Native Factory from the ground up takes longer than deploying off-the-shelf AI tools, and Siemens accepted a longer time to first visible output. The Digital Native Factory in Nanjing required a greenfield facility designed from scratch with AI at its core, a capital and governance commitment unavailable to organisations expecting AI to prove itself in existing facilities before receiving investment. The partnership with NVIDIA to build an "Industrial AI Operating System" creates a platform-level dependency: this is not a project you can walk away from if results are slower than expected. Embedding AI in workflows at the point of work, via Industrial Copilot, changes how engineers and line workers do their jobs immediately, not as a future aspiration, requiring a change management investment that incremental tool adoption does not. And deep transformation investments require board and leadership tolerance for ambiguity during the build phase. Siemens accepted this; most organisations do not have the governance structures to sustain it.

The operational results that are available are specific and credible, though full financial value attribution has not been made public. The Digital Native Factory in Nanjing demonstrated a 200% increase in manufacturing capacity and a 20% boost in productivity. At the Amberg plant, predictive maintenance AI delivered a 20% reduction in downtime. Industrial Copilot, the AI assistant embedded in industrial workflows, won the Hermes Award 2025, the industrial technology sector's leading recognition. Siemens' stated target for the broader platform is up to 50% productivity gains for industrial companies using their AI systems.

These are not marginal gains. A 200% increase in manufacturing capacity is a structural shift, not an efficiency improvement. The Hermes Award for Industrial Copilot is significant not because awards matter, but because it confirms that the interface problem, the gap between AI capability and human usability in industrial settings, was solved at a level the sector recognised as genuinely new. Siemens is currently assessed as the leading European competitor in industrial AI. That position was not inherited from 177 years of engineering heritage. It was built through deliberate, committed investment over two to three years.

The lesson is this: the right level of AI transformation is determined by your competitive position in three to five years, not by what your current processes can absorb today. If incremental AI adoption is your answer, you should be able to state clearly why that is sufficient to remain competitive as AI capability doubles every four months. If you cannot state that clearly, the question is not whether to commit more deeply, but when and at what pace. Siemens chose to answer that question in 2023. Most of its European industrial peers have not yet answered it at all.

Chapter Two

Leadership Is Not Optional

When Sponsorship Is Not Enough

Every major transformation that failed in the past twenty years had a sponsor. Someone who believed in it, championed it, hired good people, and then, gradually or suddenly, delegated it.

That is precisely how 95% of AI transformations fail.

The failure rate established in Chapter 1 is not primarily a technology problem. It is not a strategy problem. It is a leadership problem.

As Chapter 1 showed, organisations with strong executive sponsorship succeed at six times the rate of those that delegate to IT or middle management. C-level executives deeply engaged with AI are twelve times more likely to be among the top 5% of companies winning with AI.

In 2025, 72% of respondents identified the CEO as the primary AI decision-maker, up from one-third the year before. That near-doubling in a single year signals a recognition that is spreading rapidly: AI transformation is not something a leader sponsors from a distance. It is something a leader does.

This chapter makes the case for what "in the room" actually requires, for leaders, their executive teams, and their boards. The argument is not that leadership matters. The argument is that the specific kind of leadership required for this transformation is categorically different from what most leaders are currently providing, and that the gap between the two is the single most correctable cause of failure.

The distinction matters because it is correctable. Technology limitations, market timing, regulatory constraints: these are external variables that leadership can influence but not control. Whether the leader is personally, visibly, and irreversibly committed to leading the transformation is entirely within the leader's control. It is the one variable most likely to determine success, and the one most frequently abdicated.

Three Interlocking Failures

Three interlocking failures explain why delegation kills AI transformation. Each operates at a different level of the organisation, and none can be fixed in isolation.

They reinforce each other. A leader who delegates creates a board that is not challenged to govern actively. A board that does not govern actively creates space for the C-suite to fragment ownership. A fragmented C-suite creates the conditions for the micro-productivity trap.

The loop is self-sustaining, and breaking it requires intervention at all three levels simultaneously.

The delegation trap

Most leaders frame AI as an IT project, a cost centre, or an experiment to bolt onto existing processes. That framing determines everything that follows: who owns it, how it is governed, how success is measured, and when it can be defunded. It is, as we saw in Chapter 1 with the Paradigm

Shift Diagnostic, the Bolt-On response applied to a Redesign-level challenge.

The required mental model is the one Chapter 1 established: this is 80% business transformation and 20% technology transformation. That is a structural claim about where value is created and where it is lost, not a rhetorical device.

The organisations stuck in the micro-productivity trap are investing heavily in the 20% and treating the 80% as a downstream consequence. It is not. It is the primary variable.

The evidence of mismatch is stark. S&P Global reports that 42% of companies scrapped most AI initiatives in 2025, up from 17% the year before: a more than doubling of abandonment in a single year. The RAND Corporation puts AI project failure at over 80%, double non-AI IT failure rates. These are not technology failures. They are the predictable outcome of applying the wrong organisational model to a categorically different kind of change.

Meanwhile, the leaders who have adopted the right mental model are pulling away. Only 5% of firms achieve AI value at scale. Those that do commit three times the investment of average organisations and invest twice as much in workforce upskilling. They are not spending more on better algorithms. They are spending more on the organisational redesign that makes algorithms valuable.

The paradox McKinsey calls the "Gen AI Paradox" captures this precisely: nearly eight in ten companies report using generative AI, yet just as many report no significant bottom-line impact. Usage without leadership-driven architecture produces activity, not outcomes. The organisations that have broken through the paradox did so not by finding better tools but by changing who owns the transformation and how the organisation is structured to absorb it.

The board governance gap

The disconnect between AI adoption velocity and board-level governance is the single greatest enterprise risk most organisations are not measuring.

Chapter 1 set out the numbers: 88% of organisations use AI in at least one business function, yet two-thirds of board members report limited to no AI knowledge, fewer than a quarter have board-approved AI policies, and only 14% discuss AI at every meeting. Those figures have improved year-on-year, but the direction of travel is not the point. The gap itself is the point, and it has not closed.

What Chapter 1 established as a governance gap, this chapter reframes as a leadership accountability failure. The board is not underperforming because its members lack technical aptitude. It is underperforming because the CEO has not structured the governance agenda to demand better. Boards govern at the level leadership asks them to govern. When the leader frames AI as an IT project, the board asks IT-level questions.

A board that cannot distinguish between the three levels of the Paradigm Shift Diagnostic: between Bolt-On activity and genuine Redesign, cannot ask the right questions about the organisation's AI programme. A board that does not ask the right questions will approve strategies, budgets, and timelines calibrated to a different kind of change than the one actually underway.

For European boards, the governance imperative has a regulatory dimension that cannot be deferred. The EU AI Act creates direct obligations on organisational leadership: prohibited-risk provisions applied from February 2025, AI literacy training is already required, and full high-risk system requirements apply from August 2026.

Penalties reach up to 35 million euros or 7% of global annual turnover. These are not IT compliance tasks. They are leadership accountability

obligations that sit with the board and executive team.

The Act requires organisations to complete an AI inventory with risk classification, clarify their supply chain role, implement risk management systems, establish AI literacy training, and conduct fundamental rights impact assessments. Each of these is a leadership-level decision, not an IT task.

European boards that treat the EU AI Act as a legal department matter are misallocating both risk and attention. The compliance clock runs in parallel with the competitive clock, and both demand the same thing: governance that operates at the level of the transformation, not beneath it.

The C-suite alignment deficit

Transformation stalls when it is owned by one function. The pattern is consistent: a CEO announces an AI strategy, appoints a Chief Digital Officer or Chief AI Officer, funds a programme, and then returns to other priorities. The appointed leader has accountability without authority. The rest of the C-suite has neither accountability nor urgency. The programme fragments.

The gap between leadership perception and organisational reality is wider than most executives realise. Employees are using AI three times more than their leaders estimate. Ninety-one per cent of employees already use generative AI at work. Only 1% of companies report AI maturity.

Consider what that gap means in practice. Across the organisation, individuals are already building workflows, automating tasks, and experimenting with AI tools. They are doing so without coordination, without governance, and without any connection to the organisation's strategic priorities.

Individual adoption massively outpaces organisational readiness because individual adoption follows tools and convenience, whilst organisational

maturity requires deliberate architecture.

One in three workers actively sabotages corporate AI rollouts. Forty-one per cent of millennial and Gen Z employees engage in non-compliance with AI policies. Three-quarters of organisations are at or past the change saturation point. These are not signs of employee resistance to technology. They are signs of an organisation where leadership has not created the conditions for technology to be adopted coherently.

The biggest barrier to scaling AI is not employees. They are ready. The barrier is leaders who are not steering fast enough.

The role that research consistently identifies as critical: the domain leader at N-2 or N-3 level who bridges strategy and execution, exists in only 6% of companies in any meaningful form. Without these leaders, without shared accountability across every C-suite function, the leader's commitment dissipates through the organisation like heat through an uninsulated building. The intent is there. The architecture to sustain it is not.

As Allianz COO Barbara Karuth-Zelle has observed, digital and AI transformations are not mainly about the technology but rather the mindset, the people, and the organisation. That observation from a senior European executive leading one of the continent's largest AI programmes is structural, not rhetorical.

The C-suite alignment deficit is not a people problem in the human resources sense. It is a design problem. The executive team has not been structured to govern this kind of change, and most organisations have not yet restructured it.

What 'In the Room' Requires

The argument of this chapter reduces to a specific claim: the gap between what AI transformation requires of leadership and what leadership is

currently providing is the single most correctable cause of the 95% failure rate. Correcting it requires action at three governance levels simultaneously.

Leaders must lead from the front, personally and visibly. Trailblazers, the 15% of leaders research identifies as achieving AI value at scale, spend six or more hours per week upskilling on AI. They commit three times the investment of average leaders and invest twice as much in workforce development.

Half of all leaders believe their job stability depends on getting AI right. Among non-leaders, more than half believe that leaders or boards should resign if the company loses market share due to inadequate AI strategy. This is not a distant risk. It is a present expectation.

Leading from the front means more than allocating budget and attending briefings. It means reframing AI from an IT project to an operating model redesign. It means being personally fluent enough to challenge assumptions, shape priorities, and make the architectural decisions that determine whether the organisation is building at the Bolt-On level or the Redesign level.

It means signalling irreversibility: making clear, through actions and resource allocation, that this transformation is not a programme that can be quietly wound down at the next quarterly review.

The C-suite must share accountability. Each executive team member must own AI outcomes within their domain. The CFO owns the financial architecture of AI investment. The CHRO owns the workforce transformation. The COO owns the process redesign. The CTO provides the technical infrastructure.

No single role can carry the full weight. When AI is owned by one function, it is governed as a function-level initiative. When it is owned by the executive team collectively, it is governed as a transformation.

The domain leader role at N-2 or N-3 level: the person who bridges strategy and execution within each function, is what research describes as "probably the single most critical role" in scaling AI. These are not new hires. They are existing leaders given explicit accountability, development, and the authority to make AI-related decisions within their domains. Without them, the space between executive intent and operational reality remains empty.

Shared accountability also prevents the micro-productivity trap from the other direction. When each function owns its AI outcomes, use cases are designed to connect across functions rather than optimise within silos. The difference between a portfolio of disconnected AI pilots and an organisation-wide transformation is not the number of pilots. It is whether anyone has the authority and the incentive to connect them.

Boards must move from passive oversight to active governance. The minimum standard is clear: AI must be a standing agenda item at every board meeting, not a periodic update. Boards should demand specific metrics, not project-level status reports. The metrics that matter are organisational health indicators: ROI by business unit, the percentage of processes that are AI-enabled, workforce reskilling progress, regulatory alignment with the EU AI Act, and human override rates in agentic systems.

Beyond metrics, the board must determine the organisation's AI posture. Is the organisation designing at the Bolt-On level or the Redesign level? Is the current pace of transformation sufficient given the rate of competitive change? Are the governance frameworks in place to manage agentic systems that make autonomous decisions? These are board-level questions, not management-level questions, and they require board members who are fluent enough in AI to ask them with precision.

Governance maturity determines the adoption ceiling. Research shows that in high-maturity organisations, 57% of business units trust and are

ready to use AI, versus only 14% in low-maturity organisations. The board's governance posture is not a downstream consequence of the AI programme. It is a structural determinant of how far the programme can go.

Over 60% of European firms are still at the earliest stages of AI maturity. Western Europe scores well on aggregate AI readiness indices, but this masks a governance maturity gap.

The board knowledge deficit is consistent across geographies, and in Europe the accountability obligations of the EU AI Act make it acutely consequential. A board that lacks AI fluency in a jurisdiction that assigns direct legal accountability for AI governance is not merely underperforming. It is exposed.

Leadership Is the Precondition

Leadership is the precondition, not one factor among many. Without personal commitment from the leader, shared accountability across the executive team, and active governance from the board, the frameworks, playbooks, and strategies in the rest of this book will not produce results. The evidence does not leave room for ambiguity on this point.

Navrina Singh, CEO of Credo AI, put the timeline starkly: if you are not using AI as a company, you are going to be pretty irrelevant in the next 18 to 24 months. That is not hyperbole. It is a description of the competitive dynamics already visible in the data.

The argument of this chapter is settled by the evidence, not by assertion. Six times the success rate with executive-led transformations. Twelve times the likelihood of being in the top 5% when leaders are deeply engaged. A near-doubling of abandonment rates when they are not. The data points all in one direction.

The question is no longer whether to lead, but what the destination looks like. Now that the leadership precondition is established, the next chapter answers the question every committed leader asks next: what exactly are we building? Chapter 3 defines the Agentic Organisation in concrete, operational terms, introducing the AI Transformation Framework and the concept of the agentic employee.

Questions for the Board

When was the last time you personally reviewed an AI initiative's progress, not as a sponsor receiving a briefing, but as a leader shaping decisions?

Can every member of your executive team articulate how AI changes their function's operating model within the next 18 months?

Does your board receive AI-specific metrics at every meeting, or is AI a periodic agenda item?

If your organisation's AI transformation were judged solely by what you personally do each week, would it succeed?

Have you defined what "irreversible commitment" looks like for your organisation, or is your AI strategy still reversible by the next quarterly review?

Case Study: DBS Bank: The CEO Who Refused to Delegate

1. Strategic Dilemma

In the early 2010s, DBS Bank faced a choice that many large incumbents have faced and most have resolved by taking the easier path. The bank, headquartered in Singapore and one of the largest in Asia, carried a reputation summarised bluntly by its own staff: "Damn Bloody Slow."

The dilemma was not whether to transform, but how to govern the transformation. The first path was to appoint a Chief Digital Officer, fund an innovation unit, set ambitious targets, and let technology experts lead. That path offered the CEO political cover, preserved his bandwidth for other priorities, and was how most peer institutions were approaching the problem.

The second path was for the CEO to personally own the transformation, reframe the organisation's competitive identity from the top, and accept that failure would be his failure. The first path was safer in the short run. The second path was more demanding and carried greater personal risk. Piyush Gupta chose the second.

2. Decision Context

The external pressure at the time was real but not yet existential. Fintech challengers were growing rapidly across Southeast Asia, and technology giants were expanding into financial services. For a regional bank with deep retail and corporate roots, the threat was not immediate collapse but gradual commoditisation: a slow erosion of margin, relevance, and talent to faster-moving competitors.

The internal context was equally challenging. DBS was a large, complex institution with entrenched processes, multiple legacy systems, and a workforce shaped by decades of banking norms.

Previous attempts at digital improvement had produced isolated gains without systemic change, a pattern familiar to most large organisations.

The critical decision Gupta made was not simply to invest in technology. It was to redefine the organisation's competitive reference point entirely. He introduced what the bank called the GANDALF framework: a standing decision to benchmark DBS not against other banks but against Google, Amazon, Netflix, Apple, LinkedIn, and Facebook.

This was a CEO-level act of strategic reframing. It changed what "good" meant across the organisation, what hiring standards looked like, and what the board expected to see. It was irreversible in the sense that it publicly committed the organisation to standards it could not quietly walk back. That kind of commitment cannot be delegated; it can only be made by the person accountable for the organisation's direction.

3. Trade-offs

The path Gupta chose came with costs that are rarely acknowledged in the retelling:

• **Leadership bandwidth.** Personal involvement at this level is not a part-time commitment. Gupta's sustained engagement with the AI programme consumed time and political capital that could have been directed elsewhere. Other priorities were necessarily deprioritised.
• **Speed in the short term.** A delegated model, with a specialist team empowered to move fast, can show early results more quickly. CEO-led transformation moves at the pace of the CEO's own learning curve and the organisation's ability to absorb direction from the top.
• **Stakeholder management complexity.** When the CEO is visibly

accountable, every setback is visible too. There is no organisational buffer between the chief executive and the results. Gupta was publicly associated with both the ambition and the execution.

• **Pressure on the C-suite.** A CEO who personally leads a transformation raises expectations for his direct reports. Functions that might otherwise manage AI at arm's length were required to engage directly with the transformation's targets and metrics.

• **The cost of the GANDALF standard.** Benchmarking against technology companies rather than banks raised the bar for recruitment, product development, and data capability in ways that required sustained investment before they produced returns.

4. Outcome

The financial results are specific enough to stand without embellishment. The economic value DBS attributed to AI grew from S$180 million in 2022 to S$750 million in 2024, with S$1 billion projected for 2025. Return on equity doubled, from 8.4% to 18%. The share price tripled over the transformation period. These are not one-off efficiency gains; they reflect a sustained change in the bank's operating model and competitive position.

DBS is one of only two banks globally, alongside JPMorgan Chase, to have publicly disclosed total realised AI value. That transparency is itself a governance signal: it reflects a discipline of measurement and accountability that begins with what the CEO models and the board demands.

The transformation did not move from digital banking to AI-native operations in a single step; it progressed as a continuous programme, sustained by consistent leadership engagement rather than project-by-project governance.

The case is not a clean success story with no uncertainty along the way. Transforming a legacy institution to technology-company standards takes years, involves failures and course corrections, and requires the CEO to hold the line when the organisation's natural

instinct is to revert to what it knows. What the DBS case demonstrates is not that the outcome was certain, but that it was made more likely by the governance structure at the top.

5. Transferable Lesson

DBS Bank is in Singapore, not Europe, and the regulatory environment, ownership structure, and competitive context differ. The governance principles, however, are directly transferable.

What Gupta did that most European CEOs have not done is make a public, irreversible commitment, in the form of a redefined competitive reference point, that he then stayed accountable to as the transformation scaled. The transferable lesson for a European CEO is not to copy the GANDALF framework. It is to identify the equivalent act in your organisation: the decision or commitment that is visible enough, specific enough, and consequential enough that it cannot be quietly reversed at the next quarterly review.

Personal, visible, and irreversible are not adjectives. They are the governance conditions under which AI transformation succeeds. DBS is the clearest available evidence that when those conditions are present at the top, the organisation follows.

Chapter Three

What the Agentic Organisation Looks Like

You Cannot Build What You Cannot Describe

You cannot build what you cannot describe. Before the playbook, before the phased roadmap, before the first pilot, a leader needs a concrete working definition of the destination. Not a consultant's abstraction. Not a vendor's product demo. A shared, operational picture of what the organisation looks like after transformation.

Chapters 1 and 2 established two preconditions: the scale of the shift and the non-negotiable role of leadership. Most leaders now accept both. Acceptance, however, is not the same as clarity. Ask a leadership team to describe the Agentic Organisation they are building, not the tools they are buying, and the answer is typically vague, technology-centric, or borrowed wholesale from a vendor pitch.

That gap between commitment and clarity is where investment scatters, pilots multiply without connecting, and the micro-productivity trap persists. Every subsequent decision, from hiring to governance to technology investment, needs an anchor. This chapter provides it.

The Agentic Organisation is built on agentic employees. The shift from task-doer to orchestrator is not a minor adjustment in job description. It is a fundamental redefinition of what it means to be productive, skilled, and valuable. The structured model that organises the full transformation, the AI Transformation Framework, comes in Chapter 5. This chapter lays the human and conceptual foundation.

The Definition Gap

Eight in ten companies now use generative AI. Eight in ten report no significant bottom-line impact. That paradox, first identified in Chapter 1, has a specific cause: organisations lack a coherent model of what they are building towards.

The numbers bear this out. Eighty per cent of AI deployment effort is consumed not by the AI itself but by data engineering, stakeholder alignment, governance, and workflow integration. Only 12% of organisations have truly AI-ready data. Only 2% have deployed AI agents at scale, despite widespread interest and vendor noise.

Confidence in fully autonomous agents dropped from 43% to 27% in a single year. That decline does not reflect disappointment with AI capability. It reflects a growing recognition that deploying agents without redesigning how the organisation works around them produces fragile, ungovernable results.

Meanwhile, adoption continues at pace. Nearly 80% of executives report already adopting AI agents, and 88% plan budget increases. Yet most are treating AI agents as bolt-ons to existing workflows. The operating model transformation is not happening. Over 40% of agentic AI projects are predicted to fail by 2027, not because the technology underperforms, but because the organisations deploying it have not defined what success looks like beyond the pilot.

As ING's Chief Analytics Officer Bahadir Yilmaz has observed: "The conversation is no longer about if AI works. It's about understanding where it works and, more importantly, where it matters to the client." That is the European pragmatic lens at its best: the problem is not technology but where and how to deploy it.

The gap is organisational, not technical. As Chapter 1 established, algorithms account for just 10% of what is required to generate real AI value. The technology backbone is another 20%. The remaining 70% is people and processes. Without a shared picture of how those people and processes operate in the transformed state, how people, process, product, technology, and data fit together, investment scatters across disconnected pilots.

The micro-productivity trap persists because there is no shared picture of what lies beyond it. More technology does not solve this. A clear destination does.

The Agentic Organisation and the Agentic Employee

The Agentic Organisation, as defined in Chapter 1, is a company redesigned so that humans and AI agents work together as a unified system. This chapter makes that definition concrete by focusing on its central element: the people inside it.

May Habib, CEO of Writer, framed the shift precisely at the World Economic Forum in 2025: "The enterprise was designed for a world where execution was the primary constraint. Today, execution is cheap, abundant and instantaneous. Now the constraint is orchestration."

That single observation captures the transition at the heart of the Agentic Organisation. When AI agents can execute at near-zero marginal cost, the scarce resource is no longer the ability to do work. It is the ability to direct, coordinate, and judge the work that agents produce.

This is where the agentic employee enters. The term "agentic" itself means "capable of achieving outcomes independently." Applied to employees, it describes a person who is empowered, through AI, to achieve far more than they could alone, whilst maintaining accountability for outcomes. The agentic employee is not simply someone who uses AI tools. It is a fundamentally different way of working.

> ### The Agentic Employee
>
> An agentic employee is a person who works alongside autonomous AI agents as teammates, orchestrating rather than merely executing. They set goals, define boundaries, review outputs, and handle exceptions, whilst agents handle routine and complex execution. The best agentic employees create and share reusable AI workflows, producing a compounding network effect where one person's innovation multiplies the capability of everyone around them.

Three distinct profiles are emerging in organisations that have made this shift:

Profile	Role	Example
Supervisor	Orchestrates multiple AI agents across workflows	A project manager who coordinates AI agents handling scheduling, reporting, and resource allocation
Specialist	Handles exceptions and edge cases that agents cannot resolve	A compliance officer who reviews flagged transactions that agents escalated

Augmented Frontline Worker	Uses AI agents to eliminate repetitive tasks and focus on human-centric work	A customer success representative whose agent handles data gathering, letting them focus on relationship-building

These profiles are not theoretical categories. They describe roles already visible in organisations that have moved beyond bolt-on AI adoption.

The productivity evidence is substantial. A software team where output increased by 39% after AI agents became the default platform. Crucially, more experienced workers benefit more from agents, with one standard deviation higher experience correlating with 6% higher acceptance rates, whilst junior workers benefit more from tool-only AI. A legal technology trial across 12 firms found a 41% average productivity increase. Studies show time savings of 65 to 86% versus human-only workflows.

In each case, the gains came not from the AI alone but from the redesigned relationship between person and agent. The supervisor who coordinates agents effectively produces more than the specialist who works in isolation. The augmented frontline worker who shares a skill across the team multiplies impact beyond their own output.

The mindset shift required is substantial:

Old Mindset	Agentic Mindset
"I do the work"	"I orchestrate the work"
"AI is a tool I use sometimes"	"AI agents are teammates I collaborate with constantly"
"My value is what I produce"	"My value is the outcomes I enable"
"I learn skills to stay relevant"	"I create and share skills to multiply impact"
"I follow processes"	"I design and improve processes"

"I need instructions"	"I set goals and boundaries for agents, then review results"

The demand for this kind of fluency is growing at an extraordinary rate. Agentic AI fluency demand grew nearly sevenfold in two years. This is no longer a capability reserved for technical roles. It is foundational across every function: finance, HR, operations, legal, customer service, and product development all require people who can work effectively alongside agents.

The shift from left column to right in the table above is not something that happens through a training programme alone. It requires redesigned roles, new performance metrics, and a culture that rewards orchestration and skill-sharing rather than individual output. That is an organisational design challenge, not a learning and development challenge.

At the operational level, the agentic employee's work is built on two concepts that deserve explicit definition.

> ### Human-in-the-Loop
>
> Human-in-the-loop is a design principle for human-agent workflows in which human checkpoints are placed at high-stakes decision points rather than throughout every step. It ensures that agents handle routine execution autonomously whilst humans retain oversight where judgment, accountability, or ethical consideration is required.

The second concept is the building block of agentic capability.

> ### A Skill
>
> A skill is a reusable AI workflow: a discrete, repeatable sequence of agent actions built by an employee and shared across the organisation. Skills are the compound interest of the Agentic Organisation. When one

> person builds a skill and shares it, the capability multiplies across teams and functions. The best agentic employees do not just use AI; they create skills and share them, turning individual innovation into organisational capability.

Skills represent a deeper treatment that Chapter 7 will address in the context of the enterprise technology stack. For now, the critical point is that the Agentic Organisation is not built by deploying AI tools. It is built by developing people who can orchestrate agents, exercise judgment where it matters, and create reusable workflows that compound across the organisation.

The book centres on a structured model, the **AI Transformation Framework**, with four layers: Vision, Products, Operating Model, and Foundation. Chapter 5 introduces that framework in full and provides the image that ties the book's structure together. This chapter provides the human and conceptual foundation on which the entire model rests.

The connection is direct. The Operating Model layer of the framework, covering people, process, roles, and skills, is where the agentic employee lives. The Foundation layer, covering data, technology, and enablement, is where skills are built and shared. Without agentic employees, without a clear understanding of the three profiles and the mindset shift they require, the framework is an architecture without inhabitants.

The Governance Gap

Transformation at this scale creates governance demands that extend well beyond the technology function. The governance gap identified in Chapters 1 and 2, where board awareness lags behind organisational adoption, takes on a specific shape when agentic employees and AI agents are involved.

Fifty-seven per cent of organisations now have AI agents running in production. Only 21% have rigorous governance mechanisms in place. Deployment is outpacing governance by a wide margin, and the consequences are predictable: inconsistent decision boundaries, unmonitored agent actions, and the quiet proliferation of shadow agents deployed by employees without IT oversight.

Governing the agentic employee requires three elements operating together.

Tiered autonomy defines what agents and employees can do independently, with escalation to human oversight for high-impact decisions. An agent that schedules meetings operates with full autonomy. An agent that recommends credit decisions operates under supervision. An agent that triggers a financial transaction requires explicit human approval. The boundaries are risk-proportionate, not uniform.

Comprehensive logging captures every agent action for inspection, audit, and compliance. When agents make decisions at speed and scale, traceability is not optional. Without it, the organisation cannot answer the most basic governance question: what did the agent do, and on whose authority?

Circuit breakers provide real-time monitoring that can enforce boundaries and halt agent actions when they exceed defined parameters. Pre-deployment assessment, covering data permissions, system access, and risk classification, completes the picture.

Beyond these formal mechanisms, organisations must address shadow agents: the quiet proliferation of unapproved autonomous tools deployed by employees without IT oversight. When individual adoption outpaces organisational governance, which the 91% adoption versus 1% maturity gap from Chapter 1 confirms is already happening, shadow agents are not hypothetical. They are present.

As Allianz Chief Transformation Officer Maria Janssen described the approach behind Project Nemo: "AI agents support our teams by making recommendations, but the ultimate responsibility always rests with a claims professional." That principle, human accountability for agent-assisted decisions, is the governance baseline for the Agentic Organisation.

For European organisations, the governance requirements carry additional specificity. The EU AI Act creates obligations that map directly onto the dimensions of the Agentic Organisation.

AI literacy obligations under Article 4, in force since February 2025, require organisations to ensure sufficient AI literacy among all staff who interact with AI systems. Risk classification requirements, active from August 2026, demand governance of the technology and process layers.

Data protection obligations under the GDPR intersect with agentic AI in ways that are already consequential. A simple AI scheduling agent can trigger purpose creep, special-category data rules, cross-border transfer restrictions, and prohibitions on fully automated processing, all in a single workflow.

European co-determination frameworks add a governance dimension that, properly understood, becomes a design advantage rather than a constraint. Germany's Works Constitution Act requires employers to inform works councils before introducing AI systems, grants consultation rights, and triggers co-determination when AI can monitor employee behaviour. Deutsche Telekom's AI Manifesto, co-created with its Group Works Council, demonstrates that institutionalised stakeholder engagement produces faster, more durable adoption than top-down mandates.

The consequences of proceeding without proper governance are concrete. A 2025 French court suspended an AI project and imposed penalties of 1,000 euros per day for 90 days, plus 5,000 euros in damages to the works council, for deploying without consultation. Allianz, by contrast, signed the EU AI Pact in November 2024, treating proactive compliance as a

competitive differentiator rather than a cost of doing business.

Building governance into the design from the outset, rather than retrofitting it after deployment, is what separates organisations that scale from those that stall. The board needs a structured way to assess transformation progress across all dimensions rather than relying on anecdotal pilot reports. The AI Transformation Framework introduced in Chapter 5, and the governance and measurement tools developed in Chapters 12 and 13, provide that structure.

From Picture to Route

The Agentic Organisation is defined not by its technology but by its people: agentic employees who orchestrate rather than execute; three profiles that describe how work is distributed between humans and agents; skills that compound capability across the organisation; governance that keeps pace with deployment.

But a picture is not a route. The next question is: how does transformation actually scale? The answer is not "everywhere at once." It progresses through four distinct tiers, each with different governance, different leadership demands, and different measures of success. Chapter 4 introduces The Four Tiers of Transformation, from the individual to the team to the department to the organisation, and shows why skipping a tier, or assuming one tier's success automatically produces the next, is the most common structural mistake in AI programmes.

Questions for the Board

> If you asked your leadership team to describe the Agentic Organisation you are building, not the tools you are buying, would they give a consistent answer?

How many of your employees could be described as agentic employees today, and what is your plan to develop the three profiles (supervisor, specialist, augmented frontline worker) across the workforce?

Are you designing workflows around human-agent collaboration from scratch, or layering AI onto processes designed for humans working alone?

Does your board have a structured way to assess AI transformation progress across people, process, product, technology, and data, or does it rely on technology metrics and pilot updates?

What governance mechanisms are in place for employees who are already deploying AI agents without formal oversight, and how would you know if those agents made a consequential error?

Case Study: Klarna — When Automation Meets Its Limit

Strategic Dilemma

In 2023, Klarna faced a decision that most European CEOs will recognise in outline, even if the scale was unusual. The Swedish fintech was growing fast, its customer service costs were rising in proportion, and AI capabilities had reached a point where full automation of routine interactions seemed both feasible and financially compelling. CEO Sebastian Siemiatkowski drove the choice hard: deploy an AI assistant at scale, reduce headcount, and let the numbers make the argument. The alternative, a more cautious hybrid model with slower cost reduction and messier governance, looked like hesitation dressed up as prudence. Klarna chose speed.

Decision Context

The competitive logic was coherent. Klarna's AI assistant launched handling 2.3 million customer conversations in its first month alone, equivalent to 700 full-time agents. Resolution time dropped from 11 minutes to under 2 minutes. The projected annual profit improvement was $40 million. By the metrics that boards track, it was working.

Siemiatkowski was not operating in a vacuum. The broader fintech market was moving fast, and the pressure to demonstrate AI-driven efficiency to investors was real. Klarna's workforce was cut by approximately 40%. The message, internally and externally, was that the AI-first model was the future.

The problem surfaced gradually, and then clearly. Customer satisfaction for routine, low-stakes interactions remained comparable to human performance. But for complex, emotionally charged cases, the gap opened up. Klarna's own research found that 86% of customers prefer empathy over speed when their issue is genuinely difficult. No matter how fast the resolution, customers in distress

were not satisfied by a system that could not meet them where they were. Siemiatkowski later acknowledged directly: "cost unfortunately seems to have been a too predominant evaluation factor."

Trade-offs

Klarna's original decision involved real sacrifices, made under real pressure:

• **Workforce trust.** Cutting 40% of customer service staff signalled a direction that proved difficult to walk back cleanly. Re-hiring required acknowledging the reversal publicly.
• **Customer relationship quality.** Speed gains in simple cases came at the cost of satisfaction in complex ones. Those complex cases carry disproportionate weight for customer loyalty.
• **Internal governance maturity.** Moving fast meant deploying ahead of the frameworks needed to identify which interaction types were genuinely automatable and which were not.
• **Reversibility.** Once headcount is reduced and institutional knowledge dispersed, rebuilding a human service capability is slow and expensive. Klarna re-hired, but not at zero cost.
• **Leadership credibility.** A public correction of this scale requires explaining what the original decision got wrong, which is difficult when the original rationale was built on public confidence.

Outcome

Klarna partially reversed course. Human agents were re-hired for a hybrid model in which AI handles routine volume and human agents manage high-complexity, emotionally sensitive interactions. The efficiency gains from automation were real and largely retained. The experience the AI assistant delivered for straightforward queries remained strong. What changed was the design: rather than a binary choice between automation and human service, Klarna moved to a

model that assigned each type of interaction to the party best equipped to handle it.

Siemiatkowski's summary was precise: "AI gives us speed. Talent gives us empathy." That formulation matters because it reframes the question. The dilemma was never automation versus people. It was whether the organisation had a clear enough model of human-agent collaboration to design the division of work deliberately, rather than discovering it through reversal.

Klarna also continued to build outward. Its Agentic Product Protocol launched as an open standard making over 100 million products discoverable by AI agents across 12 markets. AI-driven dynamic product features with Google Cloud lifted in-app time by 15% and orders by 50%. The lesson was not that AI failed; it was that automation without deliberate design of the human layer produces a correction.

Transferable Lesson

Klarna's experience maps directly onto what this chapter argues: the Agentic Organisation is not defined by the AI it deploys, but by how it organises people, processes, products, technology, and data around human-agent collaboration. That definition requires a conscious decision about what belongs to agents and what belongs to people, made before deployment at scale rather than after. The AI Transformation Framework provides the structure for making that decision deliberately: Vision (what is the model we are building?), Products (what does the customer experience in each interaction type?), Operating Model (which employee profile handles which exception?), Foundation (what data and technology do we need to make the design work?). Klarna's reversal was not a failure of AI. It was a failure to answer those questions in order.

The decision rule for European CEOs is this: before you automate a customer-facing workflow, define which interactions require human

judgment, and build the governance to enforce that boundary. Speed and empathy are not in conflict. They require different designs.

Chapter Four

The Four Tiers of Transformation

When Individual Wins Don't Add Up

The pattern repeats across every failed AI programme. A handful of employees discover that agentic tools can compress hours of work into minutes. Their enthusiasm spreads to a team. A department runs a pilot. Results are impressive. And then nothing compounds. The organisation barely moves.

This is not a failure of ambition. It is a failure of sequence.

The first tier of transformation, the individual, is where agentic employees are made. As Chapter 3 established, these are the people who learn to work alongside AI agents as thinking partners, not search tools. Demand for AI fluency has grown sevenfold in two years. Employees who effectively use agentic tools achieve orders-of-magnitude higher output: completing tasks 25% faster, producing 40% higher quality work, and in some domains generating twice the throughput per week.

The gains are real, and they are seductive. A customer support study across more than 5,000 agents at a Fortune 500 firm found a 14% average productivity increase, with novice workers improving by 34%. Workers

with two months of AI experience performed as well as colleagues with six months of experience without AI. Manager escalations declined by 25%. Results like these generate genuine excitement, and that excitement can obscure a structural problem.

But individual fluency does not automatically become team capability. Team capability does not automatically become departmental transformation. And departmental wins do not automatically become organisational change.

Each tier requires different leadership. Each demands different governance. Each is measured by different indicators of success. The failure to recognise this is structural, not incidental. European leaders are underprepared for this transformation in part because they skip tiers, flatten them, or assume that investment at the top will cascade downward. It does not. Capability compounds only when each tier is built deliberately, and the cost of skipping a tier is not a delay. It is a collapse.

Why Capability Fails to Compound

The most striking number in recent research is the gap between individual adoption and organisational maturity. Ninety-one per cent of employees already use generative AI at work. Only 1% of companies report AI maturity. Individual adoption has massively outpaced organisational readiness.

This gap is not accidental. It is the predictable result of most leaders getting the sequence wrong: investing in enterprise-level ambitions before building the individual fluency and team-level norms on which those ambitions depend.

The evidence for bottom-up sequencing is consistent across sources. Enterprise deployments that succeed "began with power users" and empowered line managers rather than centralised AI labs. As Chapter 1

established with the BCG 10/20/70 rule, 70% of what generates real AI value is people, processes, and culture. That 70% is built tier by tier. It cannot be installed from the top.

The failure rates tell the same story. Seventy to ninety-five per cent of AI transformations fail to deliver measurable returns. Three-quarters of companies struggle to achieve value. For every 33 proofs of concept launched, only four reach production. Only 26% of AI initiatives advance beyond pilot. The consistent thread is top-down, enterprise-first approaches that bypass the foundational tiers.

The self-assessment gap compounds the problem. In one survey, 57% of leaders said they were AI-ready. Only 8% actually were. The distance between perceived readiness and real readiness is itself a tier problem: leaders who believe they have achieved Tier 2 or Tier 3 maturity when their organisations are still operating at Tier 1.

European leaders face an additional dimension. The EU AI Act, Article 4, which took effect in February 2025, makes AI literacy a legal requirement for all staff interacting with AI systems. Tier 1 is now a regulatory floor, not an optional starting point. Organisations that have underweighted individual fluency investment must now treat it as a compliance obligation as well as a strategic one.

The Four Tiers of Transformation

The sequence matters more than the ambition. A four-tier progression, Individual, Team, Department, Organisation, is the route through which AI capability actually compounds. Skipping a tier, or applying the same governance and leadership approach across all four, is the structural mistake that explains most of the failure rates cited above. The tiers are not a maturity model in the conventional sense. They are a compounding model. Each tier creates the conditions that make the next tier possible, and each

tier without its predecessor is unstable.

> The Four Tiers of Transformation
>
> The Four Tiers of Transformation describe how AI capability scales through an organisation: from Individual (personal AI fluency), to Team (shared norms and collaborative use), to Department (end-to-end workflow redesign within a function), to Organisation (cross-functional transformation and enterprise operating model). Each tier requires distinct leadership behaviour, governance, and success measures. Capability compounds only when each tier is built in sequence.

Tier 1: Individual

This is where agentic employees are made. Personal AI fluency means more than knowing how to write a prompt. It means understanding what to delegate to an AI agent, how to describe tasks effectively, how to evaluate outputs critically, and how to maintain ethical standards throughout.

Researchers have codified this as the 4D AI Fluency Framework: Delegation (knowing what to assign), Description (communicating effectively with AI), Discernment (evaluating outputs), and Diligence (maintaining ethical and safety standards). Their finding is instructive: people who treat AI as a thinking partner, engaging in augmentative conversations, exhibit more than double the fluency behaviours of those who use it transactionally.

The organisations furthest ahead have built champion programmes to accelerate Tier 1. Citi created a network of over 4,000 "AI Accelerators" across 182,000 employees in 84 countries, reaching 70% adoption. PwC Netherlands scaled from 300 AI enthusiasts to all 6,000 employees in roughly one year, using organisational network analysis to identify influential people as champions.

One of the highest-impact, lowest-cost interventions is shared prompt libraries. Organisations using structured libraries report teams becoming 40% more productive by eliminating redundant work. More telling: organisations with these libraries achieve 85% employee AI adoption, compared to 23% where usage remains purely individual.

There is, however, a critical caution. A controlled study of 758 consultants found that whilst AI users completed 12% more tasks, finished 25% faster, and produced over 40% higher quality results, those same users were 19 percentage points less likely to produce correct solutions on tasks that fell outside AI's capability frontier. Fluency without discernment is a liability. Tier 1 must build both.

Governance at Tier 1 is light but essential: personal responsibility for data handling, the ability to recognise hallucinations, and a clear tiered approach to what is permissible. Green for low-risk experimentation, yellow for tasks requiring review, red for decisions that must remain human.

Tier 2: Team

The transition from individual to team is where most transformations stall. Research shows that 61% of AI use remains at the individual level without collaborative platforms. Converting personal fluency into team capability requires something technology alone cannot provide: psychological safety.

Google's Project Aristotle study of more than 180 teams found psychological safety to be the single strongest predictor of team effectiveness, above dependability, structure, meaning, and impact. More recent research confirms this finding in the AI context specifically: 83% of executives say psychological safety measurably improves AI initiative success. Yet only 39% rate their organisation's safety as "very high," and nearly half of all workers remain uncomfortable admitting AI use to their

manager for fear of appearing incompetent or lazy.

As one leading researcher on psychological safety put it, AI adoption "should be treated as a team development effort, not just a technology upgrade." The practical marker of a Tier 2 team is a simple norm: "Have you asked the AI?" When that question becomes as natural as "Have you checked the data?", the team has crossed the threshold. Superusers emerge. Knowledge flows horizontally rather than vertically. The team begins to develop shared practices rather than individual workarounds.

The cost of failing to build this norm is measurable. Companies that grant explicit AI permissions see workers six times more likely to experiment. Trained workers are up to 19 times more likely to report productivity improvements. The permission signal is not a formality. It is a multiplier.

Structurally, Tier 2 benefits from a two-in-the-box model: a business lead and a technology lead co-owning the team's AI initiatives. This eliminates the common failure of technology being "thrown over the fence" to IT. Sun Life, Intel, and others have demonstrated that cross-functional teams of eight to twelve people, co-led in this way, are measurably more productive.

Middle managers are the critical enablers at this tier. They must become orchestrators of AI-human collaboration, yet only 35% of companies have structured change management programmes to support this transition.

Tier 3: Department

At the departmental level, transformation shifts from individual tools and team norms to end-to-end workflow redesign within a function. This is where standalone agents become multi-agent systems, and where the structural insight of agentic AI becomes visible: a human team of two to five people can supervise an "agent factory" of 50 to 100 specialised agents running an entire process.

Consider a finance workflow redesigned at Tier 3. Agent 1 extracts invoice data. Agent 2 pulls the relevant contract. Agent 3 flags discrepancies. Agent 4 drafts the resolution email. A human intervenes only at the decision point. The result: cycle times reduced by up to 80%, with improved audit trails.

In wealth management, one global firm reduced meeting preparation from two to three hours to 15 minutes by consolidating email, CRM, and portfolio data through agentic workflows. Document findability rose from 20% to 80% accessibility. Employee satisfaction increased by 22%. At Airbus, 10,000 engineers trained on AI coding tools achieved a 40% improvement in simulation cycle times. Across European organisations more broadly, industry analysis found that 56% increased profits or reduced costs through AI, with an average financial impact of over six million euros per firm.

The leap from Tier 2 to Tier 3 is where the highest proportion of initiatives die. For every 33 proofs of concept that enter production pipelines, only four emerge. The gap is not technological. It is organisational: the inability to redesign workflows, reassign authority, and restructure teams around what agentic systems make possible. Most departments add AI to existing processes rather than redesigning the process around what AI enables. The difference between those two approaches is the difference between incremental improvement and structural transformation.

Tier 4: Organisation

Cross-functional transformation is the final tier. This is where AI becomes embedded in the operating model itself, not as a tool used within functions but as the connective tissue between them.

At Tier 4, a demand surge in one region can trigger, within seconds, AI agents that reallocate advertising spend, adjust pricing, reroute stock, and

refresh creative assets across the entire organisation. An insurer can personalise campaigns across hundreds of microsegments, producing conversion rates two to three times higher and reducing call times by 25%, deployed in 16 weeks.

What makes Tier 4 possible is not superior technology. It is trust, and the organisational conditions that trust requires. Research across 1,333 executives in 27 industries found that digital-trust leaders are 1.6 times more likely to see revenue and EBIT growth of at least 10%.

A study of more than 200 large-scale transformations found that organisations captured on average only 31% of expected revenue lift and 25% of expected cost savings. The differentiator was not technology investment. Companies that invested in cultural change alongside technology saw 5.3 times higher success than those that focussed on technology alone.

The governance architecture at Tier 4 typically evolves through three phases. A centralised Centre of Excellence provides strong governance during the first 12 months. This transitions to a hub-and-spoke model, with central standards and distributed execution, between 12 and 24 months. Finally, a federated structure emerges: autonomous pods with shared governance, from 24 to 36 months onward.

Roughly 37% of large companies have established some form of CoE; the question is whether it is evolving or becoming a bottleneck.

The enterprise measure at this tier is not an AI metric. It is a business outcome: revenue growth, margin expansion, customer experience, speed to market. AI is the means. The organisation's strategic objectives are the measure.

Fewer than 10% of generative AI use cases make it past pilot to this tier. For those that do, the returns are disproportionate: the top-performing companies achieve 3.6 times the total shareholder return of their peers.

What Leaders Must Do at Each Tier

The four tiers create four distinct governance demands. Applying the same leadership approach across all of them is one of the most common and costly mistakes.

At Tier 1, the leader's role is visibility. As Chapter 2 established, leaders deeply engaged with AI are twelve times more likely to be among the top-performing companies. Only 15% qualify as "trailblazers." The single most powerful signal a leader can send at Tier 1 is personal, visible use of AI: modelling experimentation, sharing what works and what does not, and explicitly signalling that experimentation is safe.

At Tier 2, the leader's role shifts to connector. This means linking early adopters across teams, creating the conditions for horizontal knowledge flow, and keeping governance light enough that emerging norms are not crushed by premature standardisation. Middle managers require explicit support and structured change management to become the orchestrators this tier demands.

At Tier 3, the leader becomes strategist and governor. Department-level workflow redesign requires investment approval, cross-department coordination, and the willingness to let functions redesign how work is done rather than simply adding AI to existing processes. The fact that only 26% of AI initiatives advance beyond pilot reflects, in most cases, a governance failure at this tier: leaders who approve pilots but do not create the conditions for production deployment.

At Tier 4, the leader is architect. Operating model redesign, AI as corporate strategy, enterprise-wide coordination: these require active executive sponsorship. Research shows that firms with active executive sponsors are 1.8 times more likely to scale AI effectively. This is not a figurehead role. It is a design role: deciding which functions integrate, how data flows across boundaries, and where the organisation's AI-driven

operating model creates competitive advantage that individual departments cannot achieve alone.

The European advantage

European leaders often view regulation as a constraint on pace. The evidence suggests the opposite.

The EU AI Act's mandatory AI literacy requirement accelerates Tier 1 by providing budget justification and strategic urgency for the individual fluency investment that most organisations underweight. Works council consultation, far from slowing Tier 2, builds the trust and psychological safety that research identifies as the single strongest predictor of team effectiveness. GDPR's human-in-the-loop requirements enforce the governance guardrails that prevent Tier 3 and Tier 4 failures from cascading.

In Germany, works councils have co-determination rights over AI-related monitoring technology, and leading companies including Siemens and Deutsche Telekom have negotiated agile company agreements that enable AI deployment within a framework of employee trust. In France, even pilot deployments require prior works council consultation. In the Nordics, Norway has formalised a "data shop steward" role. These are not obstacles. They are structural accelerators for the trust that makes transformation compound.

Three critical mistakes

Three patterns consistently undermine tier-by-tier progress.

Over-investing in Tier 4 before Tiers 1 and 2 are in place is the first. Enterprise steering committees, platforms, and strategy are necessary, but they cannot function without a foundation of individual fluency and team-level norms. The organisation that builds the roof before the walls has an expensive structure that cannot stand.

Applying the same governance to all tiers is the second. Light-touch governance that works for individual experimentation is insufficient for departmental workflow redesign. Rigorous governance designed for enterprise deployment will suffocate early-stage adoption. Each tier requires governance calibrated to its risk profile and maturity.

Ignoring skill creation and sharing is the third. As Chapter 3 introduced, a skill is a reusable AI workflow. When individuals create skills but do not share them, the organisation loses its most powerful compounding mechanism. Shared skills are how Tier 1 fluency becomes Tier 2 capability, and how Tier 2 capability becomes Tier 3 workflow redesign.

The Route Is the Tiers

The evidence is structural, and it is consistent. Transformation scales from individuals to teams to departments to the organisation as a whole, and each tier demands a different kind of leadership. Organisations that attempt Tier 4 ambitions on Tier 1 foundations produce the failure rates that dominate the research. Organisations that build each tier deliberately, with governance and leadership calibrated to the tier, are the ones that reach the 5% who achieve AI value at scale.

The tiers are not a model to be adapted or approximated. They are a structural reality: each tier creates the preconditions for the next, and the absence of any one tier is visible in the failure rates of the tier above it.

Chapter 5 introduces the AI Transformation Framework: a four-layer map covering Vision, Products, Operating Model, and Foundation. The tiers tell you the sequence. The framework tells you the dimensions you must address at every tier. Governance, KPIs, and the Value–Capabilities–Trust decision framework are covered elsewhere in the book but are not part of the four-layer framework.

Questions for the Board: The Four Tiers

At which tier does the majority of your organisation's AI activity actually sit today: individual experimentation, team norms, departmental workflow redesign, or cross-functional enterprise transformation? What evidence do you have?

Have you invested in enterprise-level AI (steering committees, platforms, strategy) before achieving measurable individual fluency and team-level psychological safety? If yes, what is your plan to backfill the missing tiers?

What specific leadership behaviour have you or your leadership team adopted at each tier (visible use at Tier 1, connector at Tier 2, strategist at Tier 3, architect at Tier 4)?

Does your governance differ by tier, for example light-touch for individual experimentation versus rigorous for department-wide deployment, or is it one-size-fits-all?

How are you measuring success at each tier (e.g. adoption rates at Tier 1, team norms and throughput at Tier 2, workflow cycle times at Tier 3, enterprise KPIs at Tier 4)?

Case Study: ING. The Bank That Sequenced Before It Scaled

Strategic Dilemma

In the early 2020s, ING faced a choice familiar to any large European bank with genuine ambitions for AI: how to sequence the investment. Both paths had a legitimate rationale, and neither was obviously wrong.

The first path was enterprise-first: central platforms, shared data infrastructure, an AI programme announced at board level and communicated to the market. This approach would demonstrate strategic commitment, give investors a visible signal, and allow ING to match the tone of peers who were already making enterprise-scale announcements. It was the path that felt like leadership.

The second path was slower and less legible. Build tier by tier. Invest first in individual fluency and domain expertise. Establish team-level adoption norms. Redesign departmental workflows function by function before reaching for cross-functional transformation. This path would take longer to show at enterprise scale, would offer fewer memorable launch moments, and would require the organisation to govern a process rather than announce a milestone. It was the path that felt like discipline.

ING chose the second path. That choice is the reason its numbers look different from most of its peers.

Decision Context

Several factors shaped the environment in which ING made this call.

The macro backdrop was unhelpful to patience. The 2023 wave of generative AI created intense pressure on every financial institution to demonstrate AI ambition. Fintechs and digital challengers were

moving quickly. Peer banks were making announcements. The investor and analyst community was paying attention to AI strategy in a way it had not before.

Inside ING, the starting conditions were more favourable than most. The bank had already sustained serious investment in analytics capability: a 500-strong analytics team that represented years of deliberate institutional commitment rather than reactive hiring. That foundation mattered. It meant ING entered the generative AI era with a data culture already in place, domain expertise already documented, and a critical mass of people who understood what rigorous analysis looked like.

The Netherlands context added two further factors. Dutch GDPR compliance standards and the emerging EU AI Act shaped how ING approached data governance and staff obligations. And the Netherlands' strong works council culture meant that workforce engagement was not optional, it was structural. Any programme that could not withstand that scrutiny would not survive.

Bahadir Yilmaz, ING's Chief Analytics Officer, articulated the philosophy underpinning the approach plainly: "The conversation is no longer about if AI works. It's about understanding where it works and, more importantly, where it matters to the client." That orientation, client outcomes rather than capability theatre, implied a sequenced build rather than an enterprise rollout. You cannot know where AI matters to the client until you have built enough tier-by-tier capability to find out.

ING's position as the largest mortgage provider in Europe added another dimension. Mortgage origination, underwriting, and customer management involve complex domain knowledge that takes years to develop. As Yilmaz put it, "We have domain knowledge that can be transferred into digital workflows, supported by AI." That knowledge transfer only works if the individuals who hold the domain expertise are fluent enough at the individual level to participate in the redesign of the workflows that encode it.

Trade-offs

Choosing the sequenced path carried real costs.

Speed to visible enterprise results was the most immediate sacrifice. Competitors announced grander transformation programmes faster. Enterprise-first rollouts generate press releases, analyst briefings, and board presentations in a way that tier-by-tier capability building does not. ING's approach was quieter and harder to narrate externally.

Market signalling was a related cost. An enterprise AI platform, announced with appropriate fanfare, is straightforward to communicate to investors. A deliberate, function-by-function capability build is not. The story requires patience from the audience.

Governance burden was also higher in the short term. Managing a sequenced transformation means the executive team must hold a process accountable rather than mark a launch as complete. That requires more sustained attention over a longer period, without the clean milestone that a single enterprise rollout provides.

Some use cases moved more slowly than they might have in a less disciplined environment. The sequenced approach means that departmental workflow redesign waits for team-level capability to be genuinely in place, which waits for individual fluency to be genuinely in place. That sequencing has a cost in elapsed time, even when it reduces the cost of failure.

These were real trade-offs. ING made them deliberately, and the outcome data is the only honest measure of whether that discipline was justified.

Outcome

The headline number is the pilot-to-production rate: 90%, against an industry average of roughly 30%. For context, IDC found that for every 33 AI proofs of concept launched, only four reach production, a rate of approximately 12%. O'Reilly research puts the broader average at 26%. ING's rate is therefore three times the industry average and roughly seven times the IDC baseline.

That gap does not happen by accident. It reflects a structural difference in how the organisation builds. Pilots that emerge from genuine team-level adoption, within functions whose workflows have been deliberately redesigned, have a different production profile from pilots launched centrally in search of scale before the foundation is ready.

The operational outcomes reflect the same sequenced build. Seventy-five per cent of customer queries are now handled by AI chatbots, a Tier 3 outcome: end-to-end workflow redesign within the customer service function, not individual experimentation with new tools. Productivity in operations increased by 25% when AI was introduced, a result that sits at the Tier 2 and Tier 3 boundary: teams with genuine AI adoption norms embedded in redesigned workflows.

These are not outputs of a single AI launch. They are the cumulative result of building each tier before claiming the next one.

Transferable Lesson

ING's pattern maps directly onto the Four Tiers of Transformation, and the lesson for European CEOs is structural.

The 500-strong analytics team was not a Tier 4 ambition. It was Tier 1 investment sustained over time: individual capability, domain fluency, the kind of expertise that allows an employee to participate meaningfully in redesigning the workflows around them rather than simply adopting a tool they do not fully understand.

The 25% productivity gain in operations reflects Tier 2 and Tier 3 work: teams with adoption norms established, then functions with workflows genuinely redesigned. The 75% automation rate in customer service is a Tier 3 outcome, not a Tier 1 experiment that happened to work at scale.

The 90% pilot-to-production rate is the number that reveals whether an organisation has built the tiers or skipped them. When individual fluency is real, team norms are established, and departmental workflows are redesigned around what AI makes possible, pilots have somewhere to land. The organisation is ready for them.

When the tiers have been skipped and the foundation is missing, even a well-designed pilot has nowhere to go. It succeeds in isolation and stalls at the boundary of the next tier.

For a European CEO, ING's lesson is this: the governance burden and the elapsed time of sequenced transformation are not inefficiencies to be engineered away. They are the mechanism by which transformation actually compounds. The bank that sequences before it scales does not finish later. It finishes with results the enterprise-first approach rarely produces, and with a production rate that reflects an organisation genuinely ready to use what it has built.

Chapter Five

The AI Transformation Framework

Every Transformation Needs a Map

Every transformation needs a map. Without one, pilots scatter, investments fragment, and the board has no shared language for where the organisation is and what it must build. The Four Tiers from Chapter 4 and the playbook in later chapters need a common structure. That structure is the AI Transformation Framework: four layers that every organisation must address, from the top (Vision, strategy, values, objectives) to the bottom (Foundation: data, technology, enablement).

This chapter introduces the full framework and the framework image. It walks through each layer briefly so you know what each means and where the rest of the book goes deeper. No other chapter owns the full picture; this one does. Governance, KPIs, and the Value, Capabilities, Trust decision framework are essential to transformation but are covered elsewhere (Chapters 8 and 12–13). They are not part of the four-layer framework.

Why a Single Framework Matters

Scattered initiatives fail because there is no shared model of what the organisation is building. Industry surveys consistently show a gap between adoption and impact: large majorities of firms use generative AI, yet only a small fraction report significant bottom-line benefit. McKinsey's 2025 State of AI work found that only 5.5% of respondents attribute more than 5% of EBIT to AI. BCG reports that 74% of companies have yet to show tangible value, with only 4% achieving enterprise-wide AI at scale. MIT research across 300+ initiatives in 52 organisations found that 95% of generative AI pilots delivered zero measurable return.

The gap between pilot and production is where value dies. Industry analysis often shows that the vast majority of use cases that could deliver real impact remain stuck in pilot. Without a shared map, they rarely graduate to scale.

The cause is rarely technology. Deloitte notes that a large share of AI deployment effort is consumed by "unglamorous" work: data engineering, stakeholder alignment, governance, and workflow integration. Only 12% of organisations globally report truly AI-ready data; the World Economic Forum and others point to fewer than one in five organisations with high maturity in data readiness. McKinsey has found that of 25 organisational attributes tested, workflow redesign showed the biggest correlation with EBIT impact from generative AI, yet only 21% of organisations have fundamentally redesigned workflows.

The problem is the absence of a coherent, multi-dimensional model. The framework gives the board and C-suite a shared vocabulary: the four layers from Vision to Foundation. It is the basis for asking, on each layer, "Where are we, and where do we need to be?"

The AI Transformation Framework

The AI Transformation Framework is the map for the Agentic Organisation. It has four layers, top to bottom, as shown in the figure below. Each layer must be addressed; skipping one creates blind spots and limits how far transformation can scale.

Layer 1: Vision. Vision, strategy, values, and objectives sit at the top. This is the strategic direction: where the organisation is heading and why. Leadership sets it (Chapter 2). Chapter 4 connects Vision to the top of the framework and to the Four Tiers.

Layer 2: Products. This layer is organised by platform or business unit. It covers product vision and objectives, products, personas, customer journeys, and frictions. In an Agentic Organisation, products become AI-native and customer journeys dynamic and adaptive, including concepts such as zero-click discovery. Where you play and how you win in each segment depends on this layer. Later chapters go deeper by business unit.

Layer 3: Operating Model. This is where strategy becomes structure. It has two dimensions. **People:** roles, target skills and gaps, org sizing, talent journey, training, and team incentives. The agentic employee (Chapter 3) sits here: orchestrating agents, with the three profiles of supervisor, specialist, and augmented frontline worker. **Process:** org design, cross-functional teams, and product way-of-working. Transformation demands end-to-end workflow redesign around human–agent collaboration. Chapter 6 deep-dives this layer.

Layer 4: Foundation. The base of the framework has three pillars. **Data:** reusable data products, analytics, architecture, and monetisation; data as a strategic asset. **Technology:** target architecture, platforms and solutions, integration; including protocols such as MCP and agent standards. **Enablement:** culture, the AI and engineering toolbox, automation, agents, and skills. Enablement ensures that change can be driven bottom-up as well as top-down: employees understand AI, adopt an orchestrator mindset, and contribute as agentic employees; management

participates and models the behaviour. Chapter 7 covers the Foundation layer in detail.

Governance, KPIs, and the Value, Capabilities, Trust decision framework (Chapters 8 and 12–13) are essential to transformation but are not part of the four-layer framework; they are covered in later chapters.

> The AI Transformation Framework
>
> The AI Transformation Framework is the map for the Agentic Organisation. It has four layers, top to bottom: (1) Vision: vision, strategy, values, objectives; (2) Products: by platform or business unit, product vision and objectives, products, personas, customer journeys, frictions; (3) Operating Model: People (roles, target skills and gaps, org sizing, talent journey, training, team incentives), Process (org design, cross-functional teams, product way-of-working); (4) Foundation: Data (reusable data products, analytics, architecture, monetisation), Technology (target architecture, platforms and solutions, integration), Enablement (culture, AI and engineering toolbox, automation, agents, skills). Use it to describe what the Agentic Organisation is made of and how to assess readiness across each layer.

The Framework as a Governance Instrument

The framework is the basis for board and C-suite conversations. The right question on each layer is: "Where are we, and where do we need to be?" Governance demands span all four layers, not just technology. Deloitte reports that only 21% of organisations have rigorous AI governance. Research suggests that two-thirds of boards report limited to no AI knowledge and only 15% receive AI-related metrics. Relying on pilot updates or technology metrics alone is insufficient; the board should assess transformation progress against each layer.

For European organisations, EU AI Act obligations map onto the framework. Prohibitions and AI literacy requirements apply from February 2025; high-risk obligations and the majority of the Act follow from August 2026. Risk classification (unacceptable, high-risk, limited, minimal) and penalties (up to €35 million or 7% of global annual turnover) make a structured view essential.

AI literacy sits in the Operating Model; risk classification in Technology and Enablement within the Foundation; data protection in Data; product compliance in Products. The Value, Capabilities, Trust decision framework (Chapter 8) supports prioritisation and trust outcomes but is not a layer. The framework also connects to the Four Tiers (Chapter 4): each tier requires progress across the four layers; skipping a tier or a layer creates blind spots. European firms can use the framework to turn regulatory and governance obligations into a clear allocation of effort across Vision, Products, Operating Model, and Foundation.

From Map to Execution

The AI Transformation Framework provides the map. The next question is how to organise the organisation for agentic work: roles, skills, team design, and the Operating Model layer that connects vision to day-to-day execution. Chapter 6 deep-dives the Operating Model: people, process, roles, skills, and org design, and how to structure teams and roles for agentic work, from task-doer to orchestrator at scale. The framework orients; the Operating Model layer is where strategy becomes structure and where many organisations bottleneck after successful pilots.

Questions for the Board: The Framework

1. Can your leadership team name the four layers of the AI Transformation Framework (Vision, Products, Operating Model,

Foundation) and describe where your organisation stands on each?

2. Which layer is most advanced in your organisation, and which is most neglected? What is the consequence of that imbalance?

3. Does your board have a structured way to assess AI transformation progress across all four layers, or does it rely on technology metrics and pilot updates?

4. How does your AI strategy explicitly address each of the four layers, and how do you use the Value, Capabilities, Trust decision framework (Chapter 8) and governance and KPIs (Chapters 12–13) to prioritise and measure?

5. When you look at your portfolio of AI initiatives, do they map coherently onto the framework, or are they concentrated in one or two layers with gaps elsewhere?

Case Study: BBVA, Deploying One AI Transformation Map Across 120,000 Employees

1. Strategic Dilemma

BBVA faced a classic scale problem with an unusual twist: how to deploy a single, enterprise-wide AI transformation map across 120,000 employees in 25 countries without it collapsing into scattered pilots or one-size-fits-none mandates. The dilemma was not whether to adopt AI, but how to make one coherent framework the organising principle for the whole organisation. One path was to let business units and countries move at their own speed with local tools and priorities; the other was to bet on a unified vision, a central platform (ChatGPT Enterprise via a strategic alliance with OpenAI), and a shared structure that could be adapted locally. Both paths had real costs: fragmentation versus the risk of top-down rigidity and the heavy investment required to align 25 countries under one map.

2. Decision Context

BBVA had already gone through digital and mobile transformation at scale. Leadership, including Chair Carlos Torres Vila, framed the next step as entering the AI era "with even greater ambition." That meant treating AI not as a set of isolated initiatives but as a transformation that had to be mapped and governed across the full organisation. The bank operated in 25 countries with different regulatory environments, languages, and maturity levels. A single framework had to work as the shared vocabulary for the board and C-suite whilst allowing local execution. Industry evidence was sobering: the majority of use cases that could deliver real impact stay stuck in pilot, and only a small fraction of firms achieve enterprise-wide AI at scale. BBVA chose to invest in a central "AI Factory" of over 400 professionals and to embed 300 "AI wizards" in business units, tying central capability to distributed ownership.

3. Trade-offs

• **Centralised platform choice:** Committing to ChatGPT Enterprise and a strategic alliance with OpenAI meant less flexibility to mix and match best-of-breed tools by country or function; the organisation traded optionality for consistency and speed of rollout.
• **Resource concentration:** The AI Factory (400+ professionals) and 300 embedded wizards represent a large, sustained investment in central and semi-central roles; that capacity was not available for other strategic priorities.
• **Standardisation pressure:** One map across 25 countries risked under-serving local needs or forcing a single tempo where some markets were not ready; the bank accepted that some local initiatives would be delayed or reframed to fit the framework.
• **Governance and dependency:** Aligning all four layers (Vision, Products, Operating Model, Foundation) under one vision required continuous coordination and governance; the organisation took on the cost of that coordination rather than allowing siloed progress.
• **Expectation management:** Setting "even greater ambition" than digital and mobile raised the bar internally and externally; falling short would be more visible than under-promising.

4. Outcome

The outcome so far is strong adoption and early productivity signals. Around 80% of employees engage with the platform daily, and employees report saving roughly three hours per week on average. Over 3,000 custom bots have been created, indicating that the map and enablement layer are being used not only for standard workflows but for local and team-level use cases.

The structure (AI Factory plus embedded wizards) illustrates how the four layers of the AI Transformation Framework were addressed in practice: a clear vision from the top, product and journey

implications, an operating model that combines central capability with embedded roles, and a foundation that includes technology (ChatGPT Enterprise), data, and enablement (wizards, training, culture). The result is not a finished transformation but a working example of one coherent map applied at scale across a large, multinational bank.

5. Transferable Lesson

European CEOs can take one clear lesson from BBVA: the AI Transformation Framework works as a governance and organising principle only when it is used as the single map for "where we are and where we need to be" on every layer. Letting Vision, Products, Operating Model, and Foundation evolve separately leads to the pattern the chapter describes: pilots that do not scale and boards that lack a shared language.

BBVA's bet was to make the framework the basis for investment, roles, and prioritisation across 120,000 people and 25 countries, not a checklist to be filled in once. The transferable rule is to adopt one framework, use it to assess and govern all four layers, and accept the trade-offs (central platform choices, coordination cost, standardisation) in exchange for a transformation that can be measured, steered, and scaled as a whole.

Chapter Six

The Operational Layer

The Seventy Per Cent Most Organisations Ignore

Chapter 5 laid out the map: the AI Transformation Framework, four layers from Vision to Foundation. This chapter goes inside the third layer, the Operating Model, where strategy becomes structure.

BCG's 10/20/70 rule cuts to the core of why so many transformations stall. Successful AI transformation requires 70% of effort on people and processes, 20% on technology and data infrastructure, and 10% on algorithms. Most organisations get this allocation backwards. They pour budget into models, platforms, and proofs of concept whilst leaving roles, skills, and workflows unchanged.

The data confirms the gap. IBM's 2026 survey of C-suite executives found that 78% believe agentic AI requires a fundamentally new operating model. Yet only 6% of companies have begun upskilling in a meaningful way, despite 89% acknowledging the need. Nearly two-thirds of AI initiatives fail to advance beyond pilot. The problem is not technology adoption. It is operating model inertia: the organisation's roles, processes, and team structures remain designed for an era before agents.

This chapter breaks the Operating Model layer into its components. Who does the work. What skills they need. How processes run. How the organisation is structured. The answers require redesign, not incremental adjustment.

Why Pilots Succeed but Transformation Stalls

A 2026 study in the Harvard Data Science Review reported that despite 78% of companies claiming AI adoption, 80% report no measurable impact on earnings. The reason: AI is layered onto workflows designed for humans, not for human-agent collaboration.

Bain calls this the "micro-productivity trap." Grassroots experimentation with AI tools generates ideas and cultural energy, but it does not self-organise into enterprise-wide impact. Individual productivity gains stay individual. Team-level wins do not propagate. The organisation as a whole barely moves.

BCG finds that organisations which automate discrete steps within existing workflows capture only modest returns. Scaling demands end-to-end process redesign around agents as default executors. One industrial goods company redesigned its quote-to-order process using four coordinated AI agents, achieving 30 to 40% labour cost reductions and generating tens of millions in additional revenue. The difference was not the agents themselves but the process redesign that gave them room to operate.

This is the 10/20/70 insight made concrete. When an organisation invests in the 10% (algorithms) and 20% (technology) but neglects the 70% (people and process), it is structurally set up to stall. The operating model is what connects a successful pilot to scaled impact. Without it, the pilots pile up and the board loses confidence.

The Agentic Operating Model

The Operating Model layer of the AI Transformation Framework has two dimensions, as introduced in Chapter 5: **People** and **Process**. This chapter deepens both into a practical framework for redesign.

People: Three Roles the Agentic Organisation Needs

Three roles are emerging as the defining positions of the agentic era.

The Domain Leader (N-2/N-3). McKinsey describes this as "probably the single most critical role any business needs for its AI transformations." Domain leaders sit two to three levels below the CEO and own end-to-end business processes or value streams. They translate organisational vision into operational change. The gap is severe: only 17% of Fortune 500 senior leaders' skills are technical, and just 5% have held technical roles. The capability shortfall sits at the exact level where strategy becomes execution.

The Agent Orchestrator. This is the new middle-management function. An Agent Orchestrator manages synthetic AI talent: provisioning agents, connecting them to data and systems, monitoring output quality, and coordinating multi-agent workflows. BCG describes agent orchestration as the "central nervous system" of agent-first organisations, the capability that separates companies which deploy agents from those which scale them.

The AI Product Manager. This role owns the roadmap for AI agent capabilities, governs agent lifecycle, designs multi-step reasoning workflows, and establishes performance metrics. By 2026, 85% of companies expect to customise autonomous agents for specific business needs, but only one in five have mature governance models for how those agents are managed.

These roles connect to the three agentic employee profiles from Chapter 3. The Domain Leader is the executive-tier orchestrator. The Agent Orchestrator is the team-tier supervisor. The AI Product Manager governs

the interface between human intent and agent capability.

People: The Skills Shift

The shift is not from old skills to new ones. Over 70% of employer-sought skills will endure, but they will be applied differently in partnership with AI. The change is from producing first drafts to framing questions, validating outputs, and applying judgment.

AI fluency demand has grown nearly sevenfold in two years, spreading beyond technical roles into management, finance, healthcare, and frontline services. The complementary skills that matter most, critical thinking, adaptability, quality assurance, coaching, and decision-making, are human skills. They do not become less valuable; they become more so.

The most effective agentic employees create reusable AI skills and share them with colleagues, capabilities that compound as others reuse their work (the effect described in Chapter 4). Yet the gap between intent and action is stark: 89% of organisations acknowledge the need for upskilling; 6% have begun in a meaningful way. In Europe, the figures are starker still. McKinsey Global Institute estimates up to 12 million occupational transitions by 2030, double the pre-pandemic pace.

Process: Radical Workflow Redesign

Organisations must adopt what BCG calls "zero-based, outcome-driven processes": redesigning from scratch with AI at the core rather than automating legacy workflows. The returns from end-to-end redesign are measurably different from bolt-on automation: 30 to 50% business process acceleration, 25 to 40% reduction in low-value work time, and 40 to 60% cuts in engineering and design lead times. These gains materialise only when processes are redesigned as a whole.

Four multi-agent orchestration patterns are emerging. The Hierarchical pattern uses an orchestrator that breaks goals into tasks for specialist

workers, suited to complex end-to-end workflows. The Supervisor pattern assigns tasks through a coordinator whilst a separate agent evaluates quality, and dominates enterprise deployments because it mirrors human team structures. The Pipeline pattern processes work sequentially, each agent refining previous outputs, fitting linear processes such as claims handling. The Swarm pattern deploys autonomous agents in parallel with minimal coordination, suited to high-volume, low-complexity tasks. The choice of pattern depends on the process: design the workflow first, then select the pattern.

Process: Org Design

Three org design models are emerging for how organisations structure their AI capability.

Centralised. Best for low AI maturity, scarce talent, and high regulatory exposure. Strengths: standardisation, accelerated learning, clear governance. Risk: the central team becomes a bottleneck as demand scales.

Hub-and-Spoke (Federated). Best for diverse business contexts where speed matters. Strengths: business-unit ownership, local relevance. Risk: fragmentation and inconsistency without shared guardrails.

Hybrid. Combines centralised strategy, standards, and governance with embedded product teams in business units. Increasingly favoured as maturity grows. Requires a Chief AI Officer with genuine decision authority reporting to the CEO or COO, a lean AI governance council, and embedded teams owning specific outcomes. Independent risk oversight remains separate from development teams.

The typical progression: centralise first to build standards, governance, and a talent pool. Then decentralise delivery as maturity grows. The Centre of Excellence evolution (Chapter 11) maps onto this same trajectory.

> The Agentic Operating Model
>
> The Agentic Operating Model is the Operating Model layer of the AI Transformation Framework, redesigned for human-agent collaboration. It has two pillars. **People:** three new roles (Domain Leader, Agent Orchestrator, AI Product Manager) plus the skills shift from task-doer to orchestrator. **Process:** radical end-to-end workflow redesign around agents as primary actors, plus three org design patterns (Centralised, Hub-and-Spoke, Hybrid) that define how the operating model scales. Each dimension must be addressed; redesigning technology without redesigning the operating model is the primary cause of pilot purgatory.

The Operating Model as Leadership Accountability

The operating model is a board responsibility, not a delegation target. As the Allianz case study at the end of this chapter illustrates, transformation is "not mainly about the technology but rather the mindset, the people, and the organization."

Only 21% of organisations have rigorous AI governance. The gap is widest in the Operating Model layer, where decisions about roles, skills, and process redesign touch every person in the organisation. Change management must begin at the planning stage. Organisations that integrate people and process considerations from the start, co-developing solutions with frontline teams and using focused pilots with feedback loops, achieve more durable results than those that bolt on change management after the technology is deployed.

For European leaders, the operating model shift includes a dimension that does not apply in the same way elsewhere: social dialogue. Only 20% of EU trade unions currently have collective bargaining agreements addressing AI, though 42% are engaged in discussions. Workers' exposure to

algorithmic management could rise from 42% to over 55% in the medium term. This is not an obstacle. Where the operating model is redesigned with works council engagement from the outset, the result is higher trust, greater legitimacy, and more durable change. The leader who treats social dialogue as a governance advantage rather than a compliance burden builds a more resilient Agentic Organisation.

Span of control is shifting as well. McKinsey envisions one manager with a small human team orchestrating hundreds of agents running autonomously across workflows around the clock. Traditional hierarchies give way to flexible configurations of human-agent teams. This is the organisational consequence of moving from stand-alone agents through agent groups to full agentic networks, the progression mapped to the Four Tiers in Chapter 4.

From Structure to Foundation

The Operating Model layer is the 70%, and this chapter has laid out a framework for redesigning it: new roles, new skills, and new processes built for human-agent collaboration. But the operating model does not stand alone.

It depends on a foundation: data quality, technology architecture, and enablement. The two layers are tightly coupled. Process redesign requires data readiness, and 63% of organisations still lack adequate data management for AI. The skills shift requires tooling and enablement infrastructure. Org design requires platforms that enable collaboration across teams and agents. Chapter 7 dives into the Foundation layer: the final layer of the AI Transformation Framework and the base on which the operating model is built.

Questions for the Board: The Operating Model

1. Does your current operating model, roles, skills, team structures, and processes, reflect the way AI-augmented work actually gets done, or does it mirror the pre-AI organisation with AI tools layered on top?
2. Who in your organisation fills the Domain Leader role: a senior executive two to three levels below the CEO who combines business judgment with AI fluency? If no one does, what is your plan to build or hire for that capability?
3. How are you investing the 70%, people and processes, compared to the 30% for technology and algorithms? Is the ratio deliberate, or has technology absorbed most of the budget?
4. Which of the three org design models (Centralised, Hub-and-Spoke, Hybrid) describes your current AI setup, and which should it evolve towards as maturity grows?
5. Are your processes redesigned end-to-end for human-agent collaboration, or are agents layered onto legacy workflows? What would a zero-based redesign of your most important process look like?

Case Study: Allianz, Restructuring Roles Whilst Scaling Agents

1. Strategic Dilemma

Allianz, Europe's largest insurer, faced a dilemma that crystallises the Operating Model challenge. By 2025, the company had over 400 generative AI use cases in development or deployment, including Project Nemo, a seven-agent claims workflow that cut processing time by 80%. The technology worked. What was not clear was how to restructure an organisation of 160,000 people around this new reality. Cutting 1,800 roles whilst simultaneously scaling agentic AI meant Allianz had to redesign the operating model at speed, without destroying the trust, institutional knowledge, or regulatory standing built over decades. The two paths were both costly: restructure too fast and risk losing capability and social licence; restructure too slowly and let competitors set the new standard for operational efficiency.

2. Decision Context

Insurance is a trust business. Policyholders, regulators, and employees watch workforce decisions with scrutiny, and in Europe that scrutiny has institutional weight. Allianz was not a company in distress. It was a market leader choosing to transform from a position of strength, which raised rather than lowered the standard for how the transition would be managed.

COO Barbara Karuth-Zelle framed the decision explicitly: "Digital and AI transformations are not mainly about the technology but rather the mindset, the people, and the organization." This was not rhetoric. The operating model, roles, skills, team structures, and processes, was treated as the primary transformation target, with technology as the enabler.

The competitive context was pressing. ING had achieved a 90%

pilot-to-production rate through centralised AI infrastructure. BBVA had embedded 300 "AI wizards" in business units. Across European financial services, agentic AI was moving from experiment to operations. The European social dialogue framework shaped the boundaries. Only 20% of EU trade unions have collective bargaining agreements addressing AI, but 42% are engaged in formal discussions. Workforce restructuring at Allianz required works council negotiation, transparent communication, and a credible plan for affected employees.

3. Trade-offs

• **Speed versus legitimacy.** Moving quickly on role restructuring risked undermining the social dialogue process. Allianz chose to engage works councils from the planning stage, trading speed for legitimacy and durability.
• **Headcount clarity versus role ambiguity.** Announcing 1,800 role changes provided transparency but created anxiety before the new roles (orchestrator, exception handler, quality reviewer) were fully defined or staffed.
• **Short-term productivity versus long-term capability.** Project Nemo's 80% processing time reduction freed capacity immediately. The harder investment was building the skills and processes to redeploy that capacity into higher-value, judgment-intensive work.
• **Centralised deployment versus distributed adaptation.** Scaling 400+ use cases across a global insurer required central governance. Local markets needed room to adapt agent workflows to regional regulatory and operational contexts.

4. Outcome

Project Nemo delivered measurable results: claims processing time reduced by 80% through a seven-agent workflow that handles the bulk of routine cases and escalates exceptions to human reviewers.

Across the organisation, over 400 generative AI use cases moved through various stages of development and deployment.

The workforce restructuring is ongoing. Allianz has framed the 1,800 role changes not primarily as headcount reduction but as role transformation: shifting people from high-frequency, low-complexity execution to exception handling, judgment, and orchestration. The gap between leadership framing and lived experience is real. Not every affected employee will move into a redesigned role, and the company has been transparent about this limitation.

The works council process, whilst slower than a top-down mandate, has produced agreements that give the restructuring greater durability. The social dialogue requirements that some European leaders view as a constraint became a governance mechanism that forced clarity about which roles were changing, what new skills were needed, and how affected employees would be supported.

5. Transferable Lesson

The lesson from Allianz connects directly to the Agentic Operating Model. Redesigning People (roles, skills) and Process (workflows, org design) must happen simultaneously with technology deployment, not after it. A seven-agent workflow that eliminates 80% of processing time is a technology achievement. The decision to restructure 1,800 roles, engage works councils, and invest in new skills is an operating model achievement. The second is harder, slower, and more consequential.

For European leaders operating within social dialogue frameworks, the principle is concrete: design the human role and the agent workflow at the same time, with governance built in from the start. The 70% is not a follow-on project. It is the project.

Chapter Seven

Foundations of the AI Transformation

The Readiness Paradox

Ninety-eight per cent of leaders report increased urgency to deploy AI. Readiness is declining.

That is not a misprint. The Cisco AI Readiness Index, surveying over 8,000 organisations across 30 markets, found that readiness scores dropped in 2024 across five of six pillars: Infrastructure, Data, Governance, Talent, and Culture. Only 13 to 14 per cent of organisations globally qualify as fully prepared, and that figure has not moved in two years.

In Europe, the picture is worse: over 60 per cent of firms remain at the earliest readiness stages, only 23 per cent have adequate GPU capacity, 66 per cent struggle to centralise data, and just 9 per cent describe their networks as flexible.

The gap has a name. Cisco calls it AI Infrastructure Debt: the accumulation of shortcuts, deferred upgrades, and unaddressed gaps across compute, networking, data, security, and talent that compounds silently as

organisations rush to deploy.

> **AI Infrastructure Debt**
>
> AI Infrastructure Debt is the accumulation of gaps, shortcuts, and deferred upgrades in compute, networking, data management, security, and talent that compounds as organisations rush to deploy AI. It is the modern form of technical debt: silent, compounding, and value-destroying. The longer it goes unaddressed, the wider the gap between AI ambition and AI reality.

The paradox is self-reinforcing. Urgency drives premature deployment. Premature deployment drives failure. Failure drives more urgency. Nearly eight in ten companies now report using generative AI, yet just as many report no significant bottom-line impact, what McKinsey has described as the "gen AI paradox."

The way to break the cycle is to build properly. This chapter lays out what the Foundation layer requires: Data, Technology, and Enablement, the three components introduced in Chapter 5 as the base of the AI Transformation Framework.

Where Most Organisations Fail

The evidence is consistent across sources. Sixty-three per cent of organisations lack AI-ready data practices. Gartner predicts that through 2026, organisations will abandon 60 per cent of AI projects unsupported by AI-ready data. The World Economic Forum reports that fewer than one in five organisations consider themselves data-ready.

Among enterprises surveyed by Fivetran, 42 per cent say that more than half of their AI initiatives have delayed, underperformed, or failed, primarily due to data readiness.

The problem extends beyond data. Only 21 per cent of organisations have fundamentally redesigned any workflows for AI. As Chapter 6 explored, workflow redesign shows the biggest correlation with bottom-line impact, but it requires data readiness, technology readiness, and people readiness simultaneously. Most organisations still layer AI onto existing processes and wonder why pilots do not scale.

BCG's 70/10/20 rule, introduced in the previous chapter, quantifies the imbalance: 70 per cent of AI value comes from people, processes, and change management; only 10 per cent from algorithms and 20 per cent from technology infrastructure.

The implication for the Foundation layer is direct. Enablement, the third component, covering culture, skills, and the human capacity to work with AI, is where value is ultimately made or lost. Yet only 6 per cent of companies have begun meaningful upskilling, and companies are missing up to 40 per cent of AI productivity gains due to talent strategy gaps.

The problem is not technology adoption. It is the absence of foundations.

Data, Technology, and Enablement

The Foundation layer, introduced in Chapter 5 as the base of the AI Transformation Framework, has three components. Each must be addressed. None is sufficient alone.

1. Data: the strategic asset.

"AI-ready data" is a specific set of requirements, not a slogan: quality, governance, lineage, ownership, accessibility, and contextualisation. Without these, every AI project reinvents the data infrastructure, consuming months and budgets before any model training begins.

The practical unit of data readiness is the data product: a reusable dataset packaged with metadata, quality controls, and access policies.

Organisations that build data products build reusable capacity. Organisations that treat data as a project-by-project problem create waste. Research consistently shows that mature data practices deliver over three times higher return on AI investments.

As foundation models commoditise, proprietary data becomes the differentiator. Well-governed, proprietary datasets that no competitor can replicate create compounding advantages. The organisations that build them early will outperform those relying on the same publicly available data as everyone else.

A practical path exists. McKinsey outlines a 120-day data readiness approach: unified AI and data foundation in the first 30 days, governance and policy controls from day 30 to 90, and secure AI operationalisation from day 90 to 120.

The specifics vary by organisation, but the principle holds: data readiness is a focussed sprint, not a multi-year programme.

2. Technology: the architecture.

The Agentic Organisation requires a technology architecture designed for AI at scale, not for isolated pilots. The emerging consensus describes a layered stack: from infrastructure (compute, networking, storage) through data foundations and models, to skills and tools, agent frameworks and orchestration, and governance and observability at the top.

Two protocols are reshaping how agents connect to the enterprise. MCP (Model Context Protocol), described as the "USB-C for AI," provides a universal standard for connecting agents to tools and data sources. A2A (Agent2Agent) enables direct agent-to-agent communication. Together, they represent the integration standards that will determine whether agents operate across systems or remain confined to single applications.

European organisations face specific infrastructure challenges. Only 23 per cent have robust GPU capacity. Two-thirds struggle to centralise data. Just 9 per cent describe their networks as flexible enough for AI workloads.

These gaps require architectural decisions, not tactical ones. When technology investments are made project by project without a target architecture, the result is AI Infrastructure Debt: a growing tangle of incompatible systems that becomes progressively harder to unwind. Organisations classified as fully prepared design their infrastructure for future demands, not just current needs.

The choice between build, buy, and partner is strategic. Platforms range from enterprise suites to open-source frameworks, each with distinct strengths. The risk lies in deciding without reference to the overall architecture, locking the organisation into choices that optimise for one use case but constrain ten others.

3. Enablement: the human foundation.

Enablement is the component most often overlooked, and the one where the most value is at stake. BCG's 70/10/20 rule makes the point: if 70 per cent of AI value comes from people, processes, and change management, then the human capacity to absorb AI is the primary value driver, not a supporting function.

Enablement means culture, AI literacy, the engineering toolbox, automation capabilities, and the organisational capacity to absorb change. Without it, technology investments sit unused and data investments remain theoretical.

The requirement goes beyond training courses. Employees need to understand AI so that change can be driven bottom-up, not only top-down. Becoming the kind of agentic employee described in Chapter 3 requires a level of AI fluency that most organisations have not yet begun to build.

The challenge is no longer just preparing people to work with AI. It is building the systems, culture, and governance that enable humans and AI to work and learn together.

Research points to a clear threshold. Employees receiving at least five hours of structured AI training, combined with coaching, show substantially higher adoption and confidence. Below that threshold, training fails to translate into changed behaviour. Yet only 6 per cent of companies have begun meaningful upskilling.

Management must participate and model the behaviour. When leaders demonstrate AI fluency and set expectations for its use, adoption accelerates. When they delegate AI to the IT department, the signal is unmistakable: AI is optional. Enablement is not only for frontline staff.

Some organisations are already tying career progression to AI adoption, requiring regular use of AI platforms as a visible input in talent and promotion discussions. Others use generous token budgets as a talent perk, signalling investment in their people's ability to work with AI.

These are early signals, but they point in the same direction: enablement must be embedded in how the organisation hires, develops, and rewards people.

Europe faces a particular challenge in scale. Up to 12 million occupational transitions are expected by 2030, double the pre-pandemic pace. The organisations that build enablement infrastructure now will absorb that transition as a competitive advantage. Those that defer it will face a workforce skills gap that no amount of technology investment can close.

Data and Technology sit in the Foundation layer, but the process of making data usable and integrating technology into workflows belongs to the Operating Model (Chapter 6). The two layers are tightly coupled. You cannot redesign workflows without data readiness, and you cannot make

data useful without processes that consume it.

Regulation as Foundation

The EU AI Act is not an overlay on the transformation. It is a foundation input.

As Chapter 5 outlined, the Act uses a risk-based classification system with four tiers: unacceptable, high-risk, limited, and minimal risk. For the Foundation layer, the critical category is high-risk.

Any AI system used in finance, HR, critical infrastructure, health, or education must have risk management, data governance, technical documentation, human oversight, accuracy testing, and quality management systems in place by August 2026.

Each of these requirements maps directly to a Foundation component: data governance and quality (Data), documentation and robustness (Technology), human oversight and AI literacy (Enablement).

The regulatory timeline creates non-negotiable milestones. AI literacy training has been required since February 2025. Transparency obligations are active. General-purpose AI model requirements apply from August 2025. Full high-risk compliance is due by August 2026. Penalties reach €35 million or 7 per cent of global turnover.

Beyond compliance, governance creates value. ISO/IEC 42001:2023, the world's first certifiable AI management system standard, provides a practical framework covering policy, risk management, data governance, lifecycle controls, transparency, and performance evaluation. Research shows that organisations with stronger AI governance deploy AI in three additional business areas, have 28 per cent more staff using AI, and report approximately 5 per cent higher revenue growth.

Beyond the AI Act itself, specific sectors face additional regulatory layers. Financial services must comply with DORA (Digital Operational Resilience Act). Healthcare faces the Medical Device Regulation for AI diagnostics. Critical infrastructure falls under NIS2 cybersecurity requirements. For organisations in regulated sectors, the Foundation layer must account for these obligations alongside the AI Act.

For leaders, the Foundation layer must be assessed against the regulatory calendar. Which planned or active AI use cases fall under the high-risk category? Does the organisation have the Data governance, Technology documentation, and Enablement capacity in place for each?

A readiness diagnostic should map each high-risk use case to its Foundation requirements and identify the gap-closing timeline. Without that specificity, compliance becomes a scramble rather than a planned milestone.

Governance is not the enemy of speed. It is the condition for sustainable speed.

From Foundations to Decisions

The Foundation layer tells the organisation what it has: the quality of its data, the maturity of its technology, the readiness of its people.

Knowing what you have is not the same as knowing what to do with it. The next question is how to decide: which AI opportunities to pursue, which to defer, and which to reject.

Chapter 8 introduces the AI Decision Framework, balancing three angles: Value, Capabilities, and Trust. The Foundation layer feeds two of those dimensions directly. Data readiness and technology maturity determine what is possible (Capabilities). Governance readiness determines what is permissible (Trust). Together, the Foundation layer and the decision

framework give leaders the tools to move from diagnosis to action.

Chapter 11 sequences the execution: Foundation, then Activation, then Scale. Chapter 12 provides the detailed readiness diagnostic. This chapter has laid out what the foundations are and why they matter. The rest of Part II and Part IV show how to use them.

Questions for the Board: Foundations

Can your leadership team describe the current state of your Data, Technology, and Enablement foundations, not in general terms, but with specific maturity scores or assessments? If not, what is preventing that clarity?

What percentage of your data assets are genuinely AI-ready: catalogued, governed, quality-controlled, and accessible to the teams that need them? How does that compare to the benchmark of 63 per cent not ready?

Is your technology architecture designed for AI at scale, or are you accumulating AI Infrastructure Debt through tactical, project-by-project decisions? Who owns the architectural roadmap?

How many hours of structured AI training has your average employee received? Your average manager? How does that compare to the five-hour threshold for meaningful adoption improvement?

Which of your current or planned AI use cases fall under the EU AI Act's high-risk category, and do you have the data governance, documentation, and human oversight in place to comply by August 2026?

Case Study: Schneider Electric: Building AI-Ready Foundations at Enterprise Scale

1. Strategic Dilemma

Schneider Electric faced a dilemma that many long-established industrial companies recognise: how does a 200-year-old industrial company with complex global manufacturing operations build AI-ready foundations at enterprise scale?

One path was to let business units and factories deploy AI projects locally, capturing quick wins without the cost of unified data, architecture, or capability building. The other was to invest in the full Foundation layer before scaling: data products and analytics, target technology architecture, and a dedicated enablement structure (AI Hub, Chief AI Officer, company-wide capability building).

Both paths had real costs. Fragmented deployment risked pilots that never scaled and the accumulation of AI Infrastructure Debt. The foundation-first path required sustained investment, central coordination, and patience before measurable returns could materialise across 200-plus factories worldwide.

2. Decision Context

Schneider Electric operates in energy management and industrial automation across 200-plus factories and 110 distribution centres. The company had already gone through digital transformation; the next step was to treat AI as a transformation that required foundations, not just projects.

Industry evidence was sobering: Gartner predicts that through 2026, organisations will abandon 60% of AI projects unsupported by AI-ready data. McKinsey research cited by Schneider suggests industrial facilities with AI report 10–15% production increases and

4–5% EBITA gains, but only when data, technology, and human capability are in place. The company needed a structure that could work across diverse manufacturing contexts, from supply chain optimisation to yield improvement and energy management.

In November 2021, Schneider appointed Philippe Rambach as Chief AI Officer, one of the earliest CAIO appointments in European industry. Rambach reports to executive leadership and leads a dedicated global AI Hub. The Hub combines technology, process, and human enablers. By 2024, it employed over 350 AI and data science experts.

The decision was to invest in all three Foundation components: Data (industrial data products, supply chain analytics, federated architecture), Technology (predictive AI, target architecture, integration), and Enablement (AI Hub, CAIO role, company-wide capability building). The organisation chose to build the base before scaling rather than deploying projects in isolation.

3. Trade-offs

• **Centralised capability investment:** The AI Hub and 350-plus experts represent a large, sustained investment in central and semi-central roles; that capacity was not available for other strategic priorities or for accelerating local pilots.
• **Speed of deployment:** Building data products, governance, and architecture before scaling meant slower deployment of initial use cases; the organisation accepted delayed returns in exchange for a foundation that could support sustained scale.
• **Standardisation pressure:** A unified approach across 200-plus factories and 110 distribution centres risked under-serving local needs or forcing a single tempo where some sites were not ready; the company accepted that some local initiatives would be delayed or reframed to fit the foundation.

• **Governance and coordination overhead:** Aligning Data, Technology, and Enablement under one structure required continuous coordination and governance; the organisation took on the cost of that coordination rather than allowing siloed progress.
• **Expectation management:** Public targets and reported figures create visibility; falling short would be more visible than under-promising.

4. Outcome

The outcome so far is measurable progress across supply chain, manufacturing, and energy optimisation. Schneider has reported €8 million in transportation cost savings from AI-powered supply chain optimisation across 240 manufacturing facilities and 110 distribution centres, using predictive models on 200,000-plus transportation policy data points and 130,000 flow constraints; 300,000 SKUs were consolidated to 1,800 product groups.

The company's internal materials also cite a 6-day reduction in inventory days and 15% average yield improvement on specific lines; related figures (10–15% production increase) are corroborated by McKinsey research cited by Schneider.

Schneider has reported over €100 million in generated value. The company has targeted €2.0–2.5 billion in cumulative industrial productivity gains for the period 2026–2030; those targets are aspirational and not yet realised.

The structure illustrates how the Foundation layer was addressed in practice: Data (industrial data products, supply chain analytics, federated architecture), Technology (predictive AI, target architecture, integration), and Enablement (AI Hub, CAIO, 350-plus experts, company-wide capability building). The result is not a finished transformation but a working example of foundation-first investment at an industrial scale.

5. Transferable Lesson

European leaders can take one clear lesson from Schneider Electric: the Foundation layer (Data, Technology, Enablement) is not optional infrastructure. It is the condition for sustainable scale.

When an organisation invests in all three components before scaling, it avoids the pattern of pilots that never graduate to production and AI Infrastructure Debt that compounds silently. Schneider's bet was to make the Foundation layer the basis for investment, roles, and prioritisation across 200-plus factories, not a checklist to be filled in once.

The transferable rule is to assess readiness across Data, Technology, and Enablement, invest in the gaps before scaling, and accept the trade-offs (central capability cost, slower initial deployment, coordination overhead) in exchange for a transformation that can be measured, steered, and scaled as a whole.

Chapter Eight

Value, Capabilities, Trust

The Thirty-One to Seven Problem

Early AI pilots delivered impressive returns. In 2023, organisations reported an average 31% ROI on their initial AI investments. By 2025, as those projects moved from pilot to scale, that figure collapsed to 7%, below most companies' cost of capital.

The pattern has repeated across industries. Broader analysis confirms the scale of the problem: only 5% of firms achieve AI value at scale, whilst 60% achieve no material value at all. MIT research found that 95% of enterprise generative AI pilots fail to deliver sustained impact on the profit and loss statement, even as global spending on generative AI reached $30–40 billion.

That 95% figure captures something important. It measured P&L impact of pilots at relatively short timescales, not long-term scaled deployments. The statistic illustrates the severity of the decision problem, not a verdict on AI's potential.

The distinction matters. Misreading pilot-stage results as definitive leads organisations to abandon promising initiatives too early or, worse, to lose

faith in AI's potential altogether.

The root cause is not technology failure. The algorithms work. The platforms are mature. The models improve monthly. What fails is the decision: which initiatives to fund, in what order, with what capabilities in place, and with what governance to protect against risk.

Organisations pick the wrong use cases, fund them without assessing whether they can deliver, and treat governance as something to address after launch.

This chapter introduces a three-angle decision lens: Value, Capabilities, Trust. It is not a layer of the AI Transformation Framework introduced in Chapter 5. It is a cross-cutting filter that sits across the entire framework and determines where to invest, how to sequence, and what to protect.

Why Most AI Investments Disappoint

The pilot-to-scale gap is the defining problem of enterprise AI in 2025. Organisations launch dozens of proofs of concept. A few deliver modest, localised gains. Almost none compound into enterprise-wide impact. This is the micro-productivity trap identified in Chapter 1: experimentation that does not compound into enterprise-wide value.

Survey data reinforces the scale of the problem. Whilst 88% of large organisations use AI in at least one function, only 39% report enterprise-level EBIT impact. The single largest driver of EBIT capture is workflow redesign, yet only 21% of companies have done it. Nearly 70% of enterprise generative AI initiatives stall before they reach operational deployment.

Budget allocation makes the gap worse. Analysis of enterprise AI spending found that over 50% goes to sales and marketing, yet the strongest ROI appears in back-office automation and business process optimisation.

This is a prioritisation failure, not a technology failure.

Three structural errors recur. First, organisations optimise for one dimension only, usually value, and ignore capability readiness and governance risk. A high-value use case that the organisation lacks the talent, data, or process maturity to deliver is not a high-value investment. It is a high-risk one.

Second, organisations fail to establish performance baselines before implementing AI. Without a baseline, there is no way to measure whether an initiative created value or simply created activity.

Third, traditional ROI models fail for AI because the systems evolve with data and integration depth. As Chapter 1 established, this is a transformation, not a deployment. Investments measured on six-month windows will almost always disappoint. Multi-year tracking is the minimum credible measurement standard.

The common thread: organisations lack a structured way to evaluate initiatives across value, capability readiness, and governance risk simultaneously. The AI Decision Framework provides that structure.

The AI Decision Framework

Every AI initiative, from a quick-win productivity tool to an end-to-end process redesign, should be evaluated through three lenses before investment is approved: Value, Capabilities, and Trust.

> ### The AI Decision Framework
>
> The AI Decision Framework is a three-element decision lens for evaluating AI investments. Every initiative is scored on Value (does it create measurable business value?), Capabilities (can the organisation

> deliver it?), and Trust (can the organisation govern it responsibly?). It is not a layer of the AI Transformation Framework; it is a cross-cutting filter applied to every investment decision across the organisation.

The three elements interact. High-value initiatives with low capability readiness need sequencing: build the capability before scaling the investment. High-value initiatives with high governance risk need a trust-first approach: establish governance before launch.

Quick wins at the Deploy level build momentum and fund the capability and trust investment that larger initiatives require. Multiple external decision frameworks have independently converged on the same three-angle structure, reinforcing that this is not a theoretical preference but a practical necessity.

Value

Where does this initiative create measurable business value? The answer is not always obvious, and the right metric depends on the organisation's maturity.

A useful taxonomy distinguishes three types of AI value. **Deploy** covers off-the-shelf tools that deliver 10–15% productivity improvements: coding assistants, meeting summarisers, document drafters.

Reshape covers end-to-end function redesign with AI-native workflows: rethinking how procurement, customer service, or finance operates with AI at the core. **Invent** covers new business models and revenue streams that were not possible before AI.

Most organisations start with Deploy and stay there. Research shows that organisations with mature AI capabilities generate 72% of their value in core functions through Reshape and Invent, not through support-function quick wins. The progression from Deploy to Reshape to Invent is not

automatic; it requires deliberate investment in capabilities and governance at each stage.

Enterprise AI ROI operates on four layers: efficiency gains, revenue and margin improvement, risk and compliance benefits, and strategic capability building. Each layer matters at different stages of maturity.

An organisation in the early stages may focus on efficiency. An organisation redesigning end-to-end workflows should measure revenue impact and strategic capability.

The quantified evidence is encouraging for those who measure correctly. AI-powered procurement has delivered cost reductions up to 27%. Finance automation has produced 90% reductions in manual data entry, with payback within eight months. Predictive maintenance programmes report 30% cost reductions and 40% less unplanned downtime.

Two measurement disciplines are non-negotiable. Value must be measured against established baselines, not aspirations. Measurement windows must be multi-year. The pattern described in the previous section, where AI looks like a poor investment at six months, often reflects a measurement problem rather than a value problem.

Capabilities

Can the organisation deliver this initiative? The answer depends on people, data, process, and technology, roughly in that order.

As Chapter 1 established, the 10/20/70 rule puts this in proportion: 10% of the effort goes to algorithms, 20% to technology infrastructure, and 70% to people, processes, and change management. Most organisations spend on the 10% and 20% but starve the 70%.

The numbers on talent readiness are sobering. Only 6% of companies have begun upskilling meaningfully, despite 89% acknowledging the need. Only 17% of senior leaders at major corporations have technical skills; just

5% held technical roles during their careers.

The domain leader, the role Chapter 2 identified as critical at two or three levels below the chief executive, is central to capability readiness. This person translates strategic vision into operational change across a portfolio of five to fifteen interrelated use cases, bridging business context and technical possibility.

Alongside domain leaders, new roles are emerging. The Agent Orchestrator manages AI agents as operational resources, a form of middle management for synthetic talent. The AI Product Manager governs the lifecycle of AI-powered products and workflows. These roles did not exist two years ago; within eighteen months, the majority of enterprises expect to deploy AI agents that require them.

Capability assessment covers four dimensions: data readiness (quality, accessibility, architecture), talent (skills, roles, upskilling pipeline), process maturity (workflow integration, change management), and technology infrastructure (platforms, integration, scalability).

An initiative that scores high on value but low on capability is not ready for investment. It is ready for capability-building. The distinction is important: the answer is not "no," it is "not yet, and here is what we need to build first."

Only 48% of AI projects reach production today, with an average of eight months from prototype to deployment. Capability gaps are the primary reason.

Trust

Can the organisation govern this initiative responsibly? Trust is not a compliance checkbox. It encompasses three dimensions: regulatory compliance, customer and user trust, and employee trust.

Regulatory compliance in Europe means alignment with the EU AI Act, GDPR, and sector-specific regulation. The EU AI Act's full high-risk system requirements apply from August 2026, with penalties reaching €35 million or 7% of global annual turnover. Over 87% of high-risk AI cases require GDPR compliance in addition to AI Act compliance, creating layered obligations that demand integrated governance.

Customer trust requires transparency and explainability. Users need to understand when they are interacting with AI and how decisions that affect them are made. This is not merely an ethical aspiration; it is a regulatory requirement under the EU AI Act for high-risk systems.

Employee trust is the most underestimated dimension. Research reveals a 23-point gap between executive confidence in AI (84%) and employee belief in its benefits (61%). That gap threatens adoption from the inside.

Employees who do not trust AI tools, the organisation's data use policies, or its direction on AI will resist, circumvent, or ignore AI-powered workflows. In Europe, where works councils and stronger employee protections shape the landscape, closing this gap is not optional.

The evidence on trust as a leading indicator is striking. In high-maturity organisations, 57% of business units trust and actively use AI solutions. In low-maturity organisations, that figure drops to 14%. High-maturity organisations sustain AI projects for three or more years at more than twice the rate of their low-maturity peers.

As one industry leader put it: "Trust underpins it all, as systems will only ever be as autonomous as they are trustworthy."

Yet only 2% of organisations have fully operationalised responsible AI practices. Three-quarters of companies have paused AI projects at some point due to risk concerns. The gap between acknowledging trust's importance and building it into operations is one of the most significant barriers to AI value at scale.

> Tiered Autonomy
>
> Tiered autonomy is a governance principle that defines the level of independence an AI agent is permitted based on the risk and impact of its actions. Low-risk tasks (scheduling, data retrieval) operate with full autonomy. Medium-risk tasks (recommendations, initial assessments) operate under human supervision. High-risk tasks (financial transactions, regulatory decisions) require explicit human approval. The boundaries are proportionate to risk, not uniform across all AI applications.

Tiered autonomy, introduced in Chapter 3 as a design principle for agentic workflows, becomes a practical governance tool within the Trust dimension. Organisations that define clear autonomy levels for each AI initiative can scale trust incrementally: start with low-risk, fully autonomous deployments, build confidence, and expand the autonomy boundary as governance matures.

Combined with the human-in-the-loop principle defined in Chapter 3, tiered autonomy gives boards a concrete mechanism for governing AI risk without paralysing deployment.

From Scorecard to Boardroom

The AI Decision Framework is not only an operational prioritisation tool. It is a governance instrument for the board.

As Chapters 1 and 2 established, board oversight of AI remains inadequate: two-thirds of board members report limited-to-no AI knowledge, only 15% receive AI-related metrics, and nearly half say AI is not yet on the board agenda.

The gap between operational AI adoption and board-level governance creates risk that compounds silently.

The framework gives the board a structured reporting mechanism. Every AI initiative presented for approval should carry its Value, Capabilities, and Trust scores. The board sees a portfolio view: which initiatives are investment-ready across all three dimensions, which need capability-building first, and which carry trust risks that require resolution before proceeding.

Escalation criteria should be explicit. Initiatives classified as high-risk under the EU AI Act require a trust review before the board approves investment. Initiatives with high value but low capability readiness require a capability-building plan with milestones before funding. The framework turns board AI discussions from abstract strategy conversations into structured portfolio decisions.

For European organisations, trust governance is not optional. But proactive compliance is proving to be competitive advantage, not constraint. Over 230 organisations have voluntarily signed the EU AI Pact, with 63% based in the EU. More than half have gone beyond mandatory requirements.

The best boards are integrating AI governance with existing enterprise risk management rather than creating isolated oversight structures. Most boards already have risk oversight committees and risk management expertise; the task is to extend that infrastructure to cover AI, not to build a parallel governance apparatus.

The board's role is to ensure that trust is treated as an enabler of scale and adoption, not as a cost centre. When compliance investment is seen only as regulatory overhead, it gets minimised. When it is recognised as the foundation for employee adoption, customer confidence, and operational scale, it receives the investment it requires.

Trust as the European Edge

For European organisations, the Trust dimension of the AI Decision Framework is not a constraint. It is a differentiator.

The EU AI Act, GDPR, and cultural expectations around data privacy create a structural advantage that, when navigated well, accelerates adoption rather than slowing it. The European Economic and Social Committee has positioned trustworthy AI as Europe's competitive advantage, not a regulatory burden.

The numbers support this. Ninety-one per cent of European business leaders expect revenue growth in 2026 amid rising AI investment. Organisations expect to allocate 5% of annual budgets to AI by 2026, up from 3% in 2025, with more than half committing to sustained investment over a five-year horizon.

European firms that have invested in responsible AI infrastructure report that governance and ethical frameworks are among the critical enablers of scale, alongside executive sponsorship and external partnerships.

There is also a productivity dimension to the trust advantage. Research suggests that AI adoption increases labour productivity by 4% on average in European firms, driven by capital deepening rather than job displacement. That finding runs counter to the displacement fears that dominate public discourse, and it gives leaders an evidence-based narrative for building employee trust.

The organisations that master all three angles of the AI Decision Framework, creating real value, building the right capabilities, and earning trust from regulators, customers, and employees, are best placed to lead the next decade. Chapter 9 makes the full case for why Europe's distinctive regulatory and cultural context is a strategic advantage, and what leaders must do to capitalise on it.

Questions for the Board: Evaluating AI Investments

When your team proposes an AI initiative, are they required to present its expected value, capability readiness, and trust implications, or only one of the three?

Can you articulate whether your current AI portfolio is weighted towards Deploy (quick wins), Reshape (function redesign), or Invent (new models), and is that balance intentional?

Which capability gaps, people, data, process, or technology, are most likely to prevent your highest-value AI initiatives from scaling?

Has your organisation established trust governance (regulatory compliance, explainability, employee engagement) as a prerequisite for scaling AI, or is it treated as an afterthought?

If your board reviewed every active AI initiative through the lens of Value, Capabilities, and Trust today, how many would pass all three?

Case Study: Allianz — Turning Regulatory Compliance into a Competitive Weapon

1. Strategic Dilemma

In 2024, Allianz, Europe's largest insurer, faced a question that every regulated European company will eventually confront: deploy AI fast and treat compliance as a constraint to manage later, or invest upfront in trust-first governance that slows initial deployment but could become a lasting competitive advantage.

Both paths carried real costs. Speed risked regulatory exposure and reputational damage in an industry built on customer trust. Governance-first risked falling behind competitors already deploying AI across claims, underwriting, and customer service at pace.

2. Decision Context

Insurance sits squarely in the crosshairs of the EU AI Act. Claims processing, fraud detection, and risk assessment all use data and algorithms in ways that regulators classify as high-risk.

For Allianz, this was not a distant concern. The EU AI Act entered into force in August 2024, with prohibitions on unacceptable-risk AI applying from February 2025 and full high-risk requirements due by August 2026. The penalties for non-compliance reach €35 million or 7% of global annual turnover.

Allianz was not starting from zero. The company had already built one of the largest AI talent pools in the insurance industry, employing roughly 10% of the entire AI workforce across the 30 largest global insurers. It ranked second globally in the Evident AI Insurance Index, a measure of AI maturity across the sector.

The driving force behind Allianz's approach was COO Barbara

Karuth-Zelle, whose conviction that transformation is about mindset, people, and organisation, not technology, was noted in Chapter 2. That conviction shaped how Allianz approached its dilemma: rather than treating compliance as a box to tick, the leadership chose to build governance and trust into the foundation of every AI initiative.

3. Trade-offs

Choosing the trust-first path came at a cost. Allianz accepted several trade-offs:

• **Speed of deployment.** Building auditable, governance-compliant AI agents takes longer than deploying off-the-shelf tools with minimal oversight. Competitors willing to move faster on AI could capture short-term efficiency gains first.
• **Higher upfront investment.** A dedicated partnership with Anthropic to develop auditable, governance-focused AI agents required investment well beyond what simpler vendor relationships would have demanded.
• **Process overhead.** Managing close to 400 GenAI use cases through a trust-first governance lens adds review layers, documentation requirements, and decision gates that purely speed-optimised organisations avoid.
• **Change management burden.** Insisting that transformation is about mindset, people, and organisation means heavier investment in culture change, training, and organisational redesign. This is slower and harder to measure than deploying new software.

4. Outcome

The results suggest the trade-offs were worth accepting, though the full picture will take years to confirm.

Allianz now has close to 400 GenAI use cases either deployed or in

active development. The flagship initiative, Project Nemo, demonstrates what governance-first AI looks like in practice: a seven-agent workflow that processes insurance claims in under five minutes, reducing claim processing time by 80%.

This is not a pilot; it is an operational system delivering measurable value in a core business process.

On the trust dimension, Allianz was among the first signatories of the EU AI Pact in November 2024, joining over 230 organisations committing to go beyond minimum regulatory requirements. The Anthropic partnership positions Allianz to deploy AI agents that are auditable by design.

What remains uncertain deserves acknowledgement. Whether the governance investment will translate into sustained competitive advantage over insurers that took a faster, less governed approach is still an open question. The EU AI Act's full high-risk requirements do not apply until August 2026, and the regulatory landscape continues to evolve.

Allianz's bet is that being ahead of compliance, rather than scrambling to catch up, will prove cheaper and more strategically valuable over time.

5. Transferable Lesson

The Allianz case illustrates a core principle of the AI Decision Framework: Value, Capabilities, and Trust are not sequential gates to pass through one at a time. They are a compounding cycle.

Allianz invested in all three simultaneously. Value (Project Nemo's 80% efficiency gain) funded further capability investment. Capabilities (a dedicated AI workforce and a people-first transformation approach) enabled the organisation to build trust at scale.

Trust (EU AI Pact commitment, Anthropic partnership, auditable agents) enabled broader adoption across the business, which in turn generated more value.

For European CEOs evaluating their own AI initiatives through the AI Decision Framework, the lesson is concrete: do not sequence trust as the final step after value is proven and capabilities are in place.

Organisations that treat trust as an afterthought end up retrofitting governance onto systems never designed for it, at far greater cost and risk. Those that build trust into the foundation from the start create a flywheel where each element of the framework reinforces the others. In regulated industries especially, trust is not the brake on AI transformation. It is the accelerator.

Chapter Nine

Europe's Distinctive Edge

The Regulation Apology

Every European CEO I speak with apologises for their regulation. They frame it as a constraint: slower, more cautious, more expensive to navigate than their American or Asian counterparts. They are wrong.

Europe's regulatory rigour, its culture of trust, and its pragmatic approach to technology adoption are not handicaps in the race to build the Agentic Organisation. They are, in the hands of a leader who understands them, a structural competitive advantage. The real question is not whether European regulation slows you down, but whether you are using it to build something your less-regulated competitors cannot easily replicate.

The EU AI Act is not a ceiling. It is a foundation.

The investment numbers seem to settle the argument. European AI venture capital in 2024 was $12.8 billion versus $80.7 billion in the United States. US companies outspend European counterparts by 45 to 70 per cent on AI across sectors. The US share of global venture capital reached 66 per cent in 2025, its highest since the early 2000s. By any investment measure, Europe is behind.

But there is a second set of numbers that rarely features in the boardroom conversation. 53 per cent of global adults trust the EU to regulate AI effectively, compared with 37 per cent for the US and 27 per cent for China. Trust in fully autonomous AI agents has declined from 43 per cent to 27 per cent in a single year.

Over half of European generative AI users believe adoption would increase with proper government regulation. The commercial premium for demonstrably trustworthy AI keeps growing.

I call this the European Advantage Triad: Clarity, Trust, Capability. Three structural advantages that regulation and culture create for leaders willing to convert compliance into competitive edge. The triad builds on the Trust element of the AI Decision Framework (Chapter 8) and anchors it in Europe's distinctive regulatory and cultural context.

The Perception Gap

The frustration is real and it comes from the top. Siemens CEO Roland Busch called the AI Act "a key reason Europe is lagging" and described the Data Act as "toxic". SAP CEO Christian Klein wants "much less regulation" and "a united Europe". Capgemini CEO Aiman Ezzat said the EU went "too far" with AI rules.

Forty-six European technology leaders wrote to the European Commission requesting a two-year delay to the AI Act's high-risk provisions. These are not unreasonable voices. They lead companies that compete globally, and they feel the weight of regulatory complexity in their operating costs and speed to market.

Yet even the critics are hedging. Busch refused to sign the 46-CEO delay letter, arguing it did not go far enough, yet signed the EU AI Pact and now positions regulation as "driving innovation" for digital product passports.

SAP achieved ISO 42001 certification. The gap between public rhetoric and private action suggests that European leaders understand the advantage even as they complain about the cost.

The Mittelstand tells a quieter version of the same story. 40 per cent of German mid-market firms now use AI, yet average AI spending fell from 0.41 per cent to 0.35 per cent of revenue between 2024 and 2025. Only 21 per cent have an AI strategy. Adoption is rising whilst investment and strategic commitment are declining: a pattern that suggests caution, not conviction.

The honest caveat deserves stating plainly. Europe's AI investment gap is stark, and regulation alone does not close it. The gap is in capital and infrastructure, not in the quality of what European companies build when they do invest. The pattern among high-performing AI organisations consistently points to governance discipline, not spending volume, as the differentiator.

The 11 per cent of European firms classified as "AI Achievers" generate 50 per cent greater revenue growth than their peers. Their defining trait is not bigger budgets or faster deployment. It is designing AI responsibly from the start: building governance into their AI programmes from day one rather than bolting it on after deployment.

The problem is not regulation. It is the failure to convert regulation into competitive advantage.

The European Advantage Triad

The European Advantage Triad names three interlocking sources of competitive advantage that European regulation and culture create when leaders treat compliance as a strategic asset rather than a cost centre.

> The European Advantage Triad: Clarity, Trust, Capability
>
> Three reinforcing sources of competitive advantage available to organisations operating within the European regulatory and cultural environment. **Clarity** is the strategic benefit of one harmonised AI framework across 27 countries. **Trust** is the commercial premium that demonstrable compliance, consumer confidence, and workforce acceptance create. **Capability** is the operational strength that compliance processes build: data quality, governance maturity, and documentation discipline. Each element reinforces the others: clarity enables trust, trust enables capability, and capability justifies the clarity of the rules.

Each element has a clear evidence base, already visible in the companies, markets, and regulatory dynamics that define European AI today.

Clarity. The EU AI Act provides one harmonised framework across 27 member states. In the United States, over 30 states have introduced or are actively considering AI-specific legislation with conflicting definitions, thresholds, and enforcement mechanisms. Colorado, California, Texas, and Illinois each take different approaches. A European company navigating this landscape faces one rulebook. An American competitor faces dozens.

As one senior technology executive described it, the EU framework is "actually really beneficial because it provides the rules of the road on how to engage", whilst the US approach creates "a really uneven playing field".[1]

The Brussels Effect extends this clarity beyond European borders. Microsoft, Google, Amazon and over 23 other general-purpose AI providers have signed the EU's Code of Practice. UK firms are drifting "towards the highest standard of AI regulation", making the EU AI Act a de facto global baseline. For a company that masters one set of rules, the reward is a governance architecture that scales across markets.

Trust. Research shows that 53 per cent of consumers would pay a premium for AI tools guaranteeing data safety. Nearly 60 per cent want stricter safety and fairness rules. With trust in fully autonomous agents declining sharply, as documented at the start of this chapter, the commercial premium for demonstrably trustworthy AI is growing, not shrinking.

The trust advantage extends to the workforce. Employees who feel secure about AI are 2.4 times more likely to embrace it than those with job insecurity. By a margin of over two to one, workers in the US, UK, and Germany reject the growing use of AI. 70 per cent in the US believe their leaders are not truthful about AI-related job losses.

Employers remain the most trusted institution on this topic, which gives organisational leaders a powerful lever.

European social dialogue, including works councils and collective bargaining, builds the workforce acceptance that determines whether AI delivers value or remains stuck in pilot. Microsoft found that works council feedback on Copilot in Germany anticipated customer concerns before they surfaced externally. This is not resistance. It is product intelligence.

A survey of over 30,000 Europeans found that more than half of generative AI users believe adoption would increase if governments regulated AI properly. The trust deficit is not an argument against AI. It is an argument for the kind of AI that European regulation demands.

Capability. The AI Act's requirements for risk management, data governance, documentation, human oversight, and post-market monitoring mirror what high-performing AI organisations already do.

The International Association of Privacy Professionals mapped ten operational AI Act impacts that build on existing GDPR processes. These range from data protection impact assessments to fundamental rights assessments, from privacy-by-design to responsible-AI-by-design. European companies with GDPR-grade governance since 2018 have a

seven-year head start on the data quality discipline that AI demands.

Industry research shows that 60 per cent of organisations globally struggle with data quality for AI. Responsible data handling is becoming a competitive advantage, and European firms have been building that capability longer than most. ISO 42001 certification is surging, with SAP, Fastweb+Vodafone, and CM.com among early adopters.

The AI governance market is forecast to grow from approximately $300 million in 2025 to over $4.8 billion by 2034. Research suggests that structured AI governance can boost SME productivity by 27 to 133 per cent depending on sector and implementation maturity. What European companies are building now, the global market will require next.

What the AI Act Demands of Leadership Now

The timeline is not theoretical. Prohibited AI practices have been in force since February 2025. AI literacy obligations apply to all deployers. General-purpose AI transparency rules took effect in August 2025.

The full penalty regime is active: up to €35 million or 7 per cent of global turnover for prohibited practices, and €15 million or 3 per cent for other breaches.

The high-risk obligations under Annex III are legally due in August 2026. The Digital Omnibus proposal may extend this to December 2027, but it remains in the preparatory phase, Parliament is divided, and the original deadline stays binding until the Omnibus passes. Leaders who plan for the extension rather than the deadline are making a bet they cannot control.

The GPAI Code of Practice, published in July 2025, covers transparency, copyright, and safety. Its signatory taskforce held its first meeting in January 2026, with enforcement beginning August 2026.

Most member states missed the deadline for designating national competent authorities. 21 of 27 had not completed designation by August 2025. Germany approved its AI implementation act in February 2026 and designated the Bundesnetzagentur, launching an AI Service Desk for SMEs. Italy passed the first comprehensive national AI law in October 2025.

The AI Liability Directive was withdrawn in October 2025, but the Product Liability Directive, which covers AI-enabled products, requires transposition by December 2026. Liability for AI-related harm is not disappearing. It is shifting to a framework that European companies must understand.

The enforcement infrastructure is immature but building. The practical question for boards is whether they will be ready when enforcement arrives.

Board accountability is the first governance question. Who owns AI compliance in your organisation? In too many companies the answer is ambiguous, defaulting to IT or legal without clear escalation to the C-suite. Compliance accountability must sit with a named executive at C-level, with a mandate that extends beyond risk mitigation to capability development.

The EU AI Pact, with over 3,200 members and 230 pledging organisations, shows that proactive compliance is becoming a signal of governance maturity. Over 50 per cent of pledgers already exceed mandatory commitments.

Labour dialogue is a governance mechanism, not a burden. Germany's Works Council Modernisation Act gives works councils co-determination rights over AI tools. Denmark's Hilfr2 agreement is the first collective agreement regulating AI in platform work.

Only 20 per cent of European unions have AI collective bargaining agreements today, but 42 per cent are in active negotiation. For leaders, this is a window to shape these agreements before the terms are set without them.

The Nordics illustrate both the opportunity and the gap. Denmark, Finland, and Sweden lead European enterprise AI adoption at 35 to 42 per cent. Yet only 6 to 11 per cent of Nordic firms scale AI to full value, and over half invest less than €500,000 annually. Adoption without investment produces the same outcome everywhere: promising pilots that never compound.

Financial services face the most complex regulatory intersection. DORA has been mandatory since January 2025. Credit scoring and insurance pricing are classified as high-risk under the AI Act. The ECB's supervisory priorities for 2026 to 2028 centre on operational resilience.

As ECB Supervisory Board Chair Claudia Buch has stated, banks must demonstrate they do not "blindly follow AI systems' recommendations". For financial institutions, the regulatory environment goes beyond compliance: it demands proof that governance produces better decisions.

From Advantage to Action

The European edge is real but conditional. Regulation creates advantage only when companies treat compliance as capability building, not box-ticking. The investment gap will not close itself. Sovereignty programmes, including France's €109 billion AI infrastructure commitment and the Franco-German sovereign AI partnership, signal ambition but require consistent execution to deliver results.

The honest position: leaders should not celebrate regulation as inherently good. They should convert regulatory requirements into operational capabilities, including data quality, governance architecture, and workforce trust, that generate returns regardless of what competitors face. The European Advantage Triad is not automatic. It requires deliberate leadership that treats each element as a strategic programme rather than an overhead line.

The companies already doing this are the subject of the next chapter: not what Europe's edge is in theory, but what the leaders did with it in practice.

Questions for the Board: Europe's Edge

Does your organisation treat EU AI Act compliance as a cost centre or as a capability-building programme? Who owns AI compliance at board level, and how are they measured?

Can you articulate the three sources of European competitive advantage (Clarity, Trust, Capability) to your management team, and does your AI strategy explicitly build on each?

How are you engaging works councils, employee representatives, or social partners in your AI deployment? Are they adversaries or accelerators?

If a major client or procurement authority required a conformity assessment tomorrow, could you produce one? What would the process reveal about your data quality, governance, and documentation?

Are you planning for the August 2026 high-risk deadline or waiting for the Digital Omnibus extension? What is the cost of being wrong?

Case Study: ING, Governance as the Accelerator

Strategic Dilemma

Chapter 4 explored ING's tier-by-tier approach to building AI capability. This case study examines a second strategic question specific to Europe's regulatory context: how to convert governance discipline into competitive advantage.

ING could allow its business units to navigate DORA, the EU AI Act, and ECB oversight independently, accepting inconsistency in exchange for local responsiveness. Or it could centralise AI governance under a single operating structure, building one compliance architecture to serve all jurisdictions simultaneously.

COO Marnix van Stiphout described the risk of the decentralised path plainly: "bits of AI all over the place," with duplicated effort, compliance blind spots, and no way to scale what worked. The cost of the centralised path was speed, autonomy, and the cultural friction of imposing structure on teams accustomed to working independently.

Decision Context

Banking is amongst the most heavily regulated sectors in Europe. ING operates under the EU AI Act, DORA (mandatory since January 2025), ECB supervisory oversight, and national financial conduct regulations across multiple jurisdictions. Credit scoring and customer risk assessment are classified as high-risk under the AI Act's Annex III.

These obligations demanded a governance architecture that could satisfy regulators across every market simultaneously. Fragmented governance was not merely inefficient. It created direct legal and financial risk.

The ECB's supervisory priorities, as described earlier in this chapter, explicitly require banks to prove that AI governance produces better decisions, not just faster ones. For ING, this meant the governance structure had to demonstrate auditability and human oversight at every decision point.

The question had moved past whether AI worked, as the bank's analytical capability (detailed in Chapter 4) had already demonstrated. The challenge was governing deployment at scale within a regulatory environment where compliance failure carried penalties of up to 7 per cent of global turnover.

Trade-offs

ING chose centralisation: what the bank's leadership calls a "conservatively aggressive" approach. All AI development was consolidated into five domains under COO oversight. An AI Ethics Committee was established with the CEO as a member. Every use case is reviewed before deployment.

The trade-offs were real:

• **Slower initial autonomy for business units.** Teams could not launch AI initiatives independently. Every project required central review and domain allocation.
• **Concentration risk in leadership.** Placing all AI under the COO created a single point of dependency. If governance became a bottleneck, the entire pipeline would stall.
• **Cultural friction.** A large analytics community accustomed to working across the bank had to adapt to a more structured operating model.
• **Upfront investment in governance infrastructure.** Building the ethics committee, review processes, and domain structure required

time and resources before producing visible returns.

Outcome

The governance structure did not slow deployment. It accelerated it. As Chapter 4 documented, ING achieved a 90 per cent pilot-to-production rate, three times the industry average. What matters for the European advantage argument is why.

Projects that passed the governance filter were better scoped, better documented, and better aligned with regulatory requirements before reaching production. Use cases that survived the ethics committee and domain review emerged with clearer boundaries, stronger data governance, and documented decision logic.

These properties, which many organisations attempt to retrofit after launch, were built in before deployment began. The governance overhead that critics of centralisation predict turned out to be front-loaded quality investment: slower at the input stage, faster at every stage thereafter.

Customer due diligence that previously took weeks was reduced to seconds. A partnership with Google produced financial crime detection results that neither organisation could have achieved alone: two to four times more financial crime detected whilst reducing false positives by 60 per cent. These outcomes emerged from governed collaboration, not unstructured experimentation.

The shift inside ING is telling. The conversation moved from whether AI works to where it matters most to the client, and governance became the mechanism that answered that question reliably across every business line.

Transferable Lesson

ING's experience illustrates all three elements of the European

Advantage Triad working in concert. **Clarity**: one governance framework across all AI activity, removing the ambiguity that fragments effort and creates compliance risk. **Trust**: CEO participation in the ethics committee, structured review of every use case, and transparent governance that satisfies regulators, employees, and customers simultaneously.

Capability: a production rate built on the documentation discipline, data quality, and process rigour that regulatory compliance demands.

For European financial institutions in particular, where DORA, the AI Act, and ECB oversight create overlapping obligations, centralised governance is not a preference. It is a practical necessity. The alternative, allowing each business unit to interpret and comply with regulations independently, creates the fragmented compliance landscape that regulators are designed to prevent.

The lesson is direct: governance in a regulated European environment is not overhead. It is the mechanism that converts pilots into production. ING did not succeed despite regulation. It succeeded because the discipline regulation demands, when embedded in operating structure, produces faster and more reliable deployment than unstructured experimentation.

[1] Graham Abell, Vice President, Public Policy, Workday, 2025.

About the author

Stephan Hartgers-Rus is a former McKinsey strategy consultant with more than 25 years of experience helping organizations create real value from technology, long before AI became a boardroom imperative. Before becoming a consultant, Stephan studied Econometrics and Computer Science, fields that are especially relevant to the rise of AI today.

During his tenure at McKinsey, Stephan led high-stakes transformation programs before moving to senior leadership roles at digital agency Mobiquity, where he served clients worldwide. Later, he spearheaded the development of several company-wide transformation programs and ways of working for both IT innovation (establishing the foundational digital innovation framework) and, later, the AI-native product development cycle. This makes him uniquely positioned to guide organisations through the next wave of AI-driven change.

Today, through his independent advisory practice, Stephan works with CEOs, CIOs, and leadership teams navigating the transition to an AI-shaped world. His premise is simple: AI is no longer a technology question. It is a fundamental leadership question. And most organizations are not yet asking the right ones.

The Agentic Organisation is the distillation of those years in the field, the patterns he has seen, the failures he has witnessed, and the handful of organizations that got it right. It is written for leaders who want strategic clarity, not hype, and who understand that the real work of AI transformation happens in the organization, not the algorithm.

Stephan is married, has two beautiful daughters, and lives in The Hague. When he is not advising companies, you will likely find him on the water, kite-surfing.